

Information about the data:

-> The dataset was taken from PIMA Indian Diabetes Database

-> This dataset is also available in the Kaggle

-> The dataset consists of the following features:

Pregnancies : Number of times pregnant

Glucose : Plasma glucose concentration over 2 hours in an oral glucose tolerance test

Blood Pressure : Diastolic blood pressure

Skin Thickness : Triceps skin fold thickness

Insulin : 2 hours serum insulin

BMI : Body Mass Index

Diabetes Pedigree Function : Likelihood of diabetes based on family history

Age : Age of the patient

Outcome(Class label) : Class variable

: 0 IF NON-DIABETIC

: 1 IF DIABETIC

OBJECTIVE:

USING MACHINE LEARNING ANALYSIS AND MODEL TO PREDICT WHETHER A PERSON HAS DIABETES OR NOT

Importing the required libraries:

Numpy : To convert data into suitable format to feed the classification model

Pandas : To read data from CSV file and to store the data in form of Dataframe for further computations

Seaborn and Matplotlib for visualizations

Sklearn : To import the model and joblib to save the model

In [4]:

```
import numpy as np
import pandas as pd
import seaborn as s
import matplotlib.pyplot as mp
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.externals import joblib
```

-> The shape of data is 768 data points with 9 features each

-> The number of features/columns for each data point are 9

In [9]:

```
data = pd.read_csv("diabetes.csv")
print(data.shape)
print(data.columns)
print(data.ndim)
```

(768, 9)

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

2

Sample Data

In [12]:

```
print(data.head(5))
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Data Pre-Processing and Cleaning:

-> Mean imputation for the null values

-> Removing outliers

-> Checking the Correlation between the features

In [14]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose           768 non-null int64
BloodPressure     768 non-null int64
SkinThickness     768 non-null int64
Insulin           768 non-null int64
BMI               768 non-null float64
DiabetesPedigreeFunction  768 non-null float64
Age               768 non-null int64
Outcome           768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [16]:


```
corr = data.corr()
print(corr)
s.heatmap(corr, xticklabels=corr.columns, yticklabels= corr.columns)
```

	Pregnancies	Glucose	BloodPressure	
SkinThickness \				
Pregnancies	1.000000	0.129459	0.141282	-0.08167
2				
Glucose	0.129459	1.000000	0.152590	
0.057328				
BloodPressure	0.141282	0.152590	1.000000	0.20737
1				
SkinThickness	-0.081672	0.057328	0.207371	1.000000
0				
Insulin	-0.073535	0.331357	0.088933	
0.436783				
BMI	0.017683	0.221071	0.281805	
0.392573				
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928
8				
Age	0.544341	0.263514	0.239528	-
0.113970				
Outcome	0.221898	0.466581	0.065068	
0.074752				
	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	
	Age	Outcome		
Pregnancies	0.544341	0.221898		

Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x2717159f048>



Observation :

-> The brighter colors indicates more correlation between the features

-> From the table and heatmap, Glucose levels, BMI, Age, Pregnancies have more correlation with the outcome

-> Splitting the dataset to Train, Test, cross check

-> Separating the class labels from the data

-> Normalizing the data

In [123]:

```
traindata = data[150:]
testdata = data[18:150]
checkdata = data[0:18]
print(traindata.shape)
print(testdata.shape)
print(checkdata.shape)
```

```
print(checkdata.shape)
```

```
(618, 9)
(132, 9)
(18, 9)
```

In [124]:

```
train_label = np.asarray(traindata['Outcome'])
train_data = traindata.drop('Outcome', 1)
test_label = np.asarray(testdata['Outcome'])
test_data = testdata.drop('Outcome', 1)
print(train_data.shape)
print(train_label.shape)
print(test_data.shape)
print(test_label.shape)
```

```
(618, 8)
(618,)
(132, 8)
(132,)
```

In [125]:

```
mean = np.mean(train_data, axis=0)
std = np.std(train_data, axis=0)

train_data = (train_data - mean)/std

test_data = (test_data - mean)/std
print(mean)
```

```
Pregnancies          3.784790
Glucose              121.736246
BloodPressure        69.199029
SkinThickness        20.771845
Insulin              81.527508
BMI                  32.161327
DiabetesPedigreeFunction  0.475333
Age                  33.174757
dtype: float64
```

TRAINING AND EVALUATING MACHINE LEARNING MODEL

In [126]:

```
model = LogisticRegression()
print(model)
model.fit(train_data, train_label)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

Out[126]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
```

```
verbose=0, warm_start=False)
```

In [127]:

```
accuracy = model.score(test_data, test_label)
accuracy = accuracy*100
print(accuracy)
```

80.3030303030303

MODEL INTERPRETATION:

To check how different features have greater impact on the model

In [128]:

```
coef = list(model.coef_[0])
print(coef)
label = train_data.columns
print(labels)
features = pd.DataFrame()
print(len(coef))
print(len(label))
features['Features'] = label
features['Weights'] = coef
print(features)
features.sort_values(by=['Weights'], ascending= True, inplace=True)
features['positive'] = features['Weights'] > 0
features.set_index("Features", inplace=True)
features.Weights.plot(kind='barh', figsize=(11,8))
features.Weights.plot(kind = 'barh', figsize=(11,8), color =
features.positive.map({True:'green', False:'black'}))
```

```
[0.49342213322606715, 1.1564127300424993, -0.2691681507254062,
0.010539700806471045, -0.18173978760864634, 0.6729906109892765, 0.352934462
6930496, 0.0832145914265438]
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age'],
      dtype='object')
```


8


8

	Features	Weights
0	Pregnancies	0.493422
1	Glucose	1.156413
2	BloodPressure	-0.269168
3	SkinThickness	0.010540
4	Insulin	-0.181740
5	BMI	0.672991
6	DiabetesPedigreeFunction	0.352934
7	Age	0.083215

Out [128]:

<matplotlib.axes._subplots.AxesSubplot at 0x27101a7fdd8>





Observation:

-> From the above plot:

-> Glucose level, BMI, Pregnancies, DiabetesPedigreeFunction have more impact on the model

-> Insulin and Blood pressure have negative influence of prediction on the model

SAVING THE MODEL:

In [129]:

```
joblib.dump([model, mean, std], "diabetes_model")
```

Out[129]:

```
['diabetes_model']
```

Observation:

To check our model is saved properly or not, we have tested the accuracy of the model on the test data which we have done

it previously and if there is no change in the accuracy then we saved our model perfectly

In [130]:

```
loaded_model, means, stds = joblib.load('diabetes_model')
score = loaded_model.score(test_data, test_label)
acc = score*100
print("The accuracy before saving the model is : 80.30303030303")
print("The accuracy after saving the model is :", acc)
```

The accuracy before saving the model is : 80.30303030303

The accuracy after saving the model is : 80.3030303030303

MODEL PREDICTION:

We have used the test data for accuracy check and also we have the Unseen data. We have the unseen data in the checkdata

In [141]:

```
checkdata.head(5)
```

Out[141]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288

In [173]:

```
sample_data = checkdata[17:]
sample_data_features = np.asarray(sample_data.drop('Outcome',1))
sample_data_features = (sample_data_features-means)/stds
print(sample_data_features)
prediction = loaded_model.predict(sample_data_features)
print(prediction)
```

```
[[ 7.  107.  74.  0.  0.  29.6  0.254  31.  ]]
[1]
```

Observation:

The model has predicted the outcome as 1, which means that the person has diabetes.

CONCLUSION:

This is the end to end data science example. Cleaning of data, pre-processing, Imputation of missing values,

Handling categorical values, dividing the data, selecting and training the model, accuracy check on the model,

saving the model, loading the model and checking accuracy with unseen data on model. This model's accuracy can

be improved with the help of domain knowledge.

