

Information about the data:

- > The dataset has been taken from sklearn.datasets
- > This is the Boston House price dataset
- > The dataset consists of 506 rows with 13 attributes

Objective:

- > Implementing Logistic Regression on the Boston House price dataset and checking the accuracy
- > Implementing Logistic Regression with Stochastic Gradient Descent and check the accuracy
- > Evaluating Logistic Regression model with and without Stochastic Gradient Descent Optimization

Importing the dataset

In [1]:

```
from sklearn.datasets import load_boston
boston_data = load_boston()
```

Information from the dataset:

- > Shape of the dataset
- > Features of the dataset
- > Dimensionality of the dataset

In [2]:

```
print(boston_data.data.shape)
print(boston_data.feature_names)
print(boston_data.data.ndim)
```

```
(506, 13)
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
2
```

The class label of the dataset were:

In [3]:

```
In [3]:
```

```
print(boston_data.target.shape)
```

```
(506,)
```

```
In [4]:
```

```
print(boston_data.DESCR)
```

```
Boston House Prices dataset
=====
```

```
Notes
```

```
-----
```

```
Data Set Characteristics:
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive
```

```
:Median Value (attribute 14) is usually the target
```

```
:Attribute Information (in order):
```

```
  - CRIM      per capita crime rate by town
  - ZN        proportion of residential land zoned for lots over
25,000 sq.ft.
  - INDUS     proportion of non-retail business acres per town
  - CHAS      Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
  - NOX       nitric oxides concentration (parts per 10 million)
  - RM        average number of rooms per dwelling
  - AGE       proportion of owner-occupied units built prior to 1940
  - DIS       weighted distances to five Boston employment centres
  - RAD        index of accessibility to radial highways
  - TAX       full-value property-tax rate per $10,000
  - PTRATIO   pupil-teacher ratio by town
  - B         1000(Bk - 0.63)^2 where Bk is the proportion of blacks by
town
  - LSTAT     % lower status of the population
  - MEDV      Median value of owner-occupied homes in $1000's
```

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
http://archive.ics.uci.edu/ml/datasets/Housing
```

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression

not address regression
problems.

****References****

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>)

In [5]:

```
import pandas as pd
df = pd.DataFrame(boston_data.data)
df.head(5)
```

Out[5]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [6]:

```
df["price"] = boston_data.target
```

In [7]:

```
df.head(5)
```

Out[7]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

Separating the Dependent feature and Independent feature

In [8]:

```
x = df.drop('price',axis=1)
```

```
y = df['price']
```

In [9]:

```
print(x.shape)
print(y.shape)
```

```
(506, 13)
(506,)
```

In [10]:

```
from sklearn.cross_validation import train_test_split
```

C:\Users\Anil Chowdary\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)

In [11]:

```
xtrain,xtest,ytrain,ytest =
train_test_split(x,y,test_size=0.25,random_state=1)
```

In [12]:

```
print(xtrain.shape)
print(xtest.shape)
print(ytrain.shape)
print(ytest.shape)
```

```
(379, 13)
(127, 13)
(379,)
(127,)
```

In [13]:

```
from sklearn.linear_model import LinearRegression
```

In [613]:

```
lr = LinearRegression()
```

In [614]:

```
lr.fit(xtrain,ytrain)
```

Out[614]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,
normalize=False)
```

In [615]:

```
pred = lr.predict(xtest)
```

In [616]:

```
from sklearn.metrics import r2_score
```

```
In [617]:
```

```
score = r2_score(pred,ytest)
```

```
In [618]:
```

```
score
```

```
Out[618]:
```

```
0.6213821726369473
```

Visualization between the Actual values and Predicted values

```
In [20]:
```

```
import matplotlib.pyplot as mp
mp.scatter(ytest,pred)
mp.xlabel('ytest')
mp.ylabel('pred')
mp.title('Actual vs Predicted')
```

```
Out[20]:
```

```
Text(0.5,1,'Actual vs Predicted')
```


Distribution of Errors:

```
In [21]:
```

```
ydelta = ytest - pred
import seaborn as s
s.set_style("whitegrid")
s.kdeplot(ydelta,bw=0.5)
```

```
Out[21]:
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x177fcce0208>
```



Distribution of Predicted prices and Actual prices

```
In [22]:
```



```

mp.figure(211)
s.kdeplot(pred,bw=0.5)
mp.xlabel("Prices")
mp.figure(212)
s.kdeplot(ytest,bw=0.5)
mp.xlabel("Prices")

```

Out[22]:

Text(0.5,0,'Prices')



Observation:

- > The actual values have spread between 3 and 44
- > The predicted values have spread between 3 and 53
- > The error distribution have more negative skewness, this means that the model is predicting more than the actual values
- > I have used R-Squared as the performance measure which resulted in 62.1 percent

In [784]:

```
print(xtrain.shape)
print(xtest.shape)
print(ytrain.shape)
print(ytest.shape)
```

```
(379, 13)
(127, 13)
(379,)
(127,)
```

Scaling the features

In [785]:

```
scale = StandardScaler()
```

In [786]:

```
xtrain = scale.fit_transform(xtrain)
xtest = scale.fit_transform(xtest)
```

In [787]:

```
from sklearn.linear_model import SGDRegressor
```

In [788]:

```
from sklearn import linear_model
```

SGD with learning rate 1 and number of iterations = 1

In [866]:

```
sgdr = SGDRegressor(alpha=1,n_iter=1,loss='squared_loss')
sgdr.fit(xtrain,ytrain)
pred = sgdr.predict(xtest)
score = r2_score(pred,ytest)*100
print(score)
```

```
-289.8165364388144
```

```
C:\Users\Anil Chowdary\Anaconda3\lib\site-
packages\sklearn\linear_model\stochastic_gradient.py:117:
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be remo-
ved in 0.21. Use max_iter and tol instead.
  DeprecationWarning)
```

SGD with learning rate 0.5 and number of iterations = 2

In [867]:

```
sgdr = SGDRegressor(alpha=0.5,n_iter=2,random_state=1,loss='squared_loss')
sgdr.fit(xtrain,ytrain)
pred = sgdr.predict(xtest)
score = r2_score(pred,ytest)*100
print(score)
```

-18.481088609809902

```
C:\Users\Anil Chowdary\Anaconda3\lib\site-  
packages\sklearn\linear_model\stochastic_gradient.py:117:  
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be remo  
ved in 0.21. Use max_iter and tol instead.  
DeprecationWarning)
```

SGD with learning rate 0.25 and number of iterations = 3

In [868]:

```
sgdr = SGDRegressor(alpha=0.25,n_iter=3,random_state=2,loss='squared_loss')  
sgdr.fit(xtrain,ytrain)  
pred = sgdr.predict(xtest)  
score = r2_score(pred,ytest)*100  
print(score)
```

28.60078816790762

```
C:\Users\Anil Chowdary\Anaconda3\lib\site-  
packages\sklearn\linear_model\stochastic_gradient.py:117:  
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be remo  
ved in 0.21. Use max_iter and tol instead.  
DeprecationWarning)
```

SGD with learning rate 0.12 and number of iterations = 4

In [869]:

```
sgdr = SGDRegressor(alpha=0.12,n_iter=4,random_state=3,loss='squared_loss')  
sgdr.fit(xtrain,ytrain)  
pred = sgdr.predict(xtest)  
score = r2_score(pred,ytest)*100  
print(score)
```

55.000206266037296

```
C:\Users\Anil Chowdary\Anaconda3\lib\site-  
packages\sklearn\linear_model\stochastic_gradient.py:117:  
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be remo  
ved in 0.21. Use max_iter and tol instead.  
DeprecationWarning)
```

In []:

SGD **with** learning rate 0.06 **and** number of iterations = 5

In [870]:

```
sgdr = SGDRegressor(alpha=0.06,n_iter=5,random_state=4,loss='squared_loss')  
sgdr.fit(xtrain,ytrain)  
pred = sgdr.predict(xtest)  
score = r2_score(pred,ytest)*100  
print(score)
```

53.57293449043019

```
C:\Users\Anil Chowdary\Anaconda3\lib\site-  
packages\sklearn\linear_model\stochastic_gradient.py:117:  
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be remo  
ved in 0.21. Use max_iter and tol instead.
```


ved in 0.21. Use max_iter and tol instead.
DeprecationWarning)

SGD with learning rate 0.03 and number of iterations = 6

In [871]:

```
sgdr = SGDRegressor(alpha=0.03,n_iter=6,random_state=5,loss='squared_loss')
sgdr.fit(xtrain,ytrain)
pred = sgdr.predict(xtest)
score = r2_score(pred,ytest)*100
print(score)
```

62.04614151659265

C:\Users\Anil Chowdary\Anaconda3\lib\site-packages\sklearn\linear_model\stochastic_gradient.py:117:
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be removed in 0.21. Use max_iter and tol instead.
DeprecationWarning)

SGD with learning rate 0.02 and number of iterations = 7

In [872]:

```
sgdr = SGDRegressor(alpha=0.02,n_iter=7,random_state=6,loss='squared_loss')
sgdr.fit(xtrain,ytrain)
pred = sgdr.predict(xtest)
score = r2_score(pred,ytest)*100
print(score)
```

64.54798575686078

C:\Users\Anil Chowdary\Anaconda3\lib\site-packages\sklearn\linear_model\stochastic_gradient.py:117:
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be removed in 0.21. Use max_iter and tol instead.
DeprecationWarning)

SGD with learning rate 0.01 and number of iterations = 8

In [873]:

```
sgdr = SGDRegressor(alpha=0.01,n_iter=8,random_state=7,loss='squared_loss')
sgdr.fit(xtrain,ytrain)
pred = sgdr.predict(xtest)
score = r2_score(pred,ytest)*100
print(score)
```

59.081650022913855

C:\Users\Anil Chowdary\Anaconda3\lib\site-packages\sklearn\linear_model\stochastic_gradient.py:117:
DeprecationWarning: n_iter parameter is deprecated in 0.19 and will be removed in 0.21. Use max_iter and tol instead.
DeprecationWarning)

Observation:

By trying different combinations of Learning rates and Number of It

rations we can see that by increasing learning rate from 0.02 there is decreasing in the performance score

In [833]:

```
print(pred.shape)
print(ytest.shape)
```

```
(127,)
(127,)
```

In [834]:

```
print(type(ytest))
print(type(pred))
```

```
<class 'pandas.core.series.Series'>
<class 'numpy.ndarray'>
```

Mean difference between the Actual price and Predicted price

In [836]:

```
print(ytest.mean())
print(pred.mean())
```

```
23.09448818897638
22.22302364393628
```

In [837]:

```
t = pd.DataFrame(data=pred)
```

Median difference between Actual price and Predicted price

In [840]:

```
print(ytest.median())
print(t.median())
```

```
22.0
0      22.417535
dtype: float64
```

In [841]:

```
print(ytest.std())
print(pred.std())
```

```
9.992229576951791
7.421183114166382
```

In [842]:

```
import numpy as np
```

In [843]:

```
print(np.sum(ytest))
print(np.sum(pred))
```

2933.0

2822.3240027799075

Weight of the features from the Model

In [844]:

```
sgdr.coef_
```

Out[844]:

```
array([-0.76933981,  0.70170576, -0.2440426 ,  0.68822587, -1.30061069,
        2.44441636, -0.16806442, -2.07190516,  0.89617887, -0.31317425,
       -1.8610315 ,  0.6094922 , -3.57303014])
```

-> Error is the difference between the Actual prices and Predicted prices

In [852]:

```
error = ytest-pred
print(error.max())
print(error.min())
print(np.sum(error))
```

16.91771118510735

-7.495529208285134

110.67599722009287

Observation:

The maximum error is 16.9177

The minimum error is -7.45


The sum of all the errors is 110.67599

In [853]:

```
s.kdeplot(error,bw=0.5)
```

Out[853]:

<matplotlib.axes._subplots.AxesSubplot at 0x177fe62d208>





In [854]:

```
sgdr.intercept_
```

Out[854]:

```
array([22.22302364])
```


Distribution of Errors

In [855]:

```
import seaborn as s
s.set_style("whitegrid")
s.kdeplot(ydelta,bw=0.5)
```

Out[855]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x177fe5c39b0>
```



Distribution of Predicted prices and Actual prices


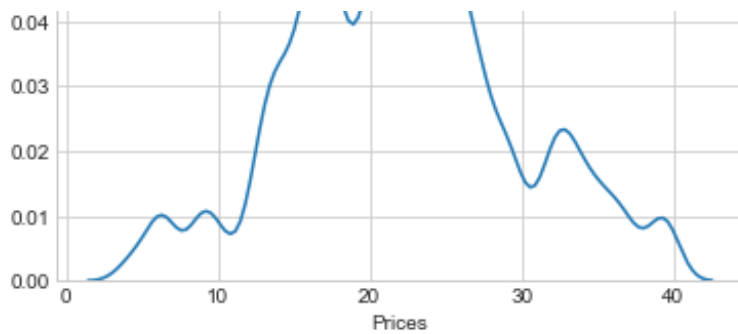
In [859]:

```
mp.figure(211)
s.kdeplot(pred,bw=1)
mp.xlabel("Prices")
mp.title("Predicted prices")
mp.figure(212)
s.kdeplot(ytest,bw=1)
mp.xlabel("Prices")
mp.title("Actual prices")
```

Out[859]:

```
Text(0.5,1,'Actual prices')
```





Observation:

Both the predicted prices and actual prices have more spread in the region of 10 and 30

Conclusion:

- > The best performance measure for regression model is R-Squared
- > The value of R-Square in the range of 0 to 1 is a good model
- > If the R-Square value is equal to zero then the model is simple mean model
- > If R-Square value is negative then the model is performing very bad
- > By using simple Linear Regression we got R-Square value as 0.62
- > By using Stochastic Gradient Descent with different combination of learning rates and

number of iterations we got the best R-Square value as 0.645