

Information about data:

->We have the amazon reviews dataset from kaggle

->Reviews are given for the product

->The features of the data were:

Id

ProductId- unique identifier for the product

UserId- unique identifier for the user

ProfileName

HelpfulnessNumerator- number of users who found the review helpful

HelpfulnessDenominator- number of users who indicated whether they found the review helpful or not

Score-rating between 1 and 5

Time-timestamp for the review

Summary- brief summary of the review

Text- text of the review

-> Based on the score of the review we classify them into positive and negative

Number of reviews: 568,454



objective:

-> Cleaning the dataset by classifying them into positive and negative reviews based on the rating provided and removing the duplicates

-> Converting the text data to vectors by using Bag of words, Tfidf, word2vec, Average word2vec

-> Applying k-fold cross-validation to determine the best k value has

ving accuracy as the check measure

-> Here we can use both random based splitting of data or time based splitting of data

-> Since it is the data of reviews they may change overtime due to modifications of products

-> The accuracy which we obtained by random based splitting may change for the future data

-> We can assure about the accuracy obtained in time base splitting for unseen future data also.

-> In this for the first method i am going to do both random based splitting and time based splitting

-> Random based splitting is possible for every dataset, but for time based splitting there should time attribute

Importing the required libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as mp
import seaborn as s
from sklearn.feature_extraction.text import
TfidfTransformer, TfidfVectorizer, CountVectorizer
import sqlite3
import nltk
import string
from sklearn.metrics import accuracy_score, roc_curve, auc, confusion_matrix
from sklearn.cross_validation import train_test_split, cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn import cross_validation
from collections import Counter
```

```
/Users/vthumati/anaconda3/lib/python3.6/site-
packages/sklearn/cross_validation.py:41: DeprecationWarning: This module wa
s deprecated in version 0.18 in favor of the model_selection module into wh
ich all the refactored classes and functions are moved. Also note that the
interface of the new CV iterators are different from that of this module. T
his module will be removed in 0.20.
```

```
"This module will be removed in 0.20.", DeprecationWarning)
```

->loading the data and information about the data

-> Shape of the data

-> Dimensionality of the data

-> Dimensionality of the data

-> Attributes of the data

In [2]:

```
con = sqlite3.connect("database.sqlite")
data = pd.read_sql_query("SELECT * FROM Reviews WHERE Score != 3", con)
print(data.shape)
print(data.ndim)
print(data.columns)

(525814, 10)
2
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
       'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')
```

Removing the Duplicates from the data

In [3]:

```
#####function to categorise rating into positive and negatives
def change(n):
    if n>3:
        return 'positive'
    return 'negative'

rating = data['Score']
#####take the ratings
rating = rating.map(change)
#####apply function change on ratings column
data['Score'] = rating
#####updating the column with positive and negatives
data.head(6)
##### head with first 6 elements in data
```

Out [3]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulness
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres	1	1

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulness
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0
5	6	B006K2ZZ7K	ADT0SRK1MGOEU	Twoapennything	0	0

Data Cleaning:Removing Duplicates

In [4]:

```
user = pd.read_sql_query("""SELECT * FROM Reviews WHERE UserId= "AR5J8UI46C
URR" ORDER BY ProductId """,con)
print(user)
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS

```

2 LOACKER QUADRATINI VANILLA WAFERS
3 LOACKER QUADRATINI VANILLA WAFERS
4 LOACKER QUADRATINI VANILLA WAFERS

```

```

                                Text
0 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

Observation:

-> Here we can see that for the same time span we got five reviews, practically which is not possible

->This happened because when the user given review for a product it is applied to all the flavors of the product

In [5]:

```
sorteddata = data.sort_values('ProductId',axis=0,ascending=True,inplace=False,kind='quicksort',na_position='last')
```

In [6]:

```
finaldata = sorteddata.drop_duplicates(subset={"UserId","ProfileName","Time","Text"},keep='first',inplace=False)
```

Information about the modified data:

-> Shape of the data

-> Dimensionality of the data

-> Attributes of the data

-> Sample of modified data

In [7]:

```
print(finaldata.shape)
print(finaldata.ndim)
print(finaldata.columns)
print(finaldata.head(5))
```

```
(364173, 10)
```

```
2
```

```
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
       'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')
```

	Id	ProductId	UserId	ProfileName
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski
138688	150506	0006641040	A2IW4PEEK02R0U	Tracy

138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"
138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"
138691	150509	0006641040	A3CMRKGE0P909G	Teresa

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time
138706	0	0	positive	939340800
138688	1	1	positive	1194739200
138689	1	1	positive	1191456000
138690	1	1	positive	1076025600
138691	3	4	positive	1018396800

	Summary
138706	EVERY book is educational
138688	Love the book, miss the hard cover version
138689	chicken soup with rice months
138690	a good swingy rhythm for reading aloud
138691	A great way to learn the months

	Text
138706	this witty little book makes my son laugh at l...
138688	I grew up reading these Sendak books, and watc...
138689	This is a fun way for children to learn their ...
138690	This is a great little book to read aloud- it ...
138691	This is a book of poetry about the months of t...

Keeping in mind about the performance capability of the box, here i am taking the subset of the data

In [8]:

```
finaldata1 = finaldata.sample(n=100000)
```

Information about the sampled data:

- > Shape of the data
- > Dimensionality of the data
- > Attributes if the data
- > Sample of modified data

In [9]:

```
print(finaldata1.shape)
print(finaldata1.ndim)
print(finaldata1.columns)
print(finaldata1.head(5))
```

```
(100000, 10)
```

```
2
```

```
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
      'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')
```

	Id	ProductId	UserId	ProfileName
425529	460208	B005BRHVD6	AAU4SEQOA9NE	Ryan Barton "The Knife"

494547	534658	B0018Q1G94	A3FITKZSOILAJX	Ursula Nemeth
147114	159580	B000V0H4A8	A5D3JR6T9T53V	Boston Gal "Boston Gal"
441781	477703	B0029TECEA	A3ONWXEVE56C7G	J. Schick
276015	299126	B004VB0Q9Y	A2Y5OR28P88B49	Patricia

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time
\				
425529	1	1	positive	1305244800
494547	0	0	positive	1346803200
147114	1	1	positive	1326672000
441781	8	9	positive	1256342400
276015	0	0	positive	1346803200

	Summary \
425529	Love it.
494547	The Best Pasta
147114	Addicted!
441781	Excellent soup
276015	A lot of K-cups - good price

	Text
425529	I think that the roctane Gu is thinner than th...
494547	I love this Pasta... as good as any "normal" p...
147114	Yummy Yummy - love this snack! slight olive o...
441781	Great tasting soup! Lots of chicken and veget...
276015	After quite of bit of searching, I found this ...

BAG OF WORDS:

-> RANDOM BASED SPLITTING:

->SIMPLE CROSS-VALIDATION AND K-FOLD CROSS-VALIDATION:

-> TO FIND THE BEST K-VALUE WITH ACCURANCY AS MEASURE

Using Count-Vectorizer to convert the text to vector representation

In [10]:

```
cv = CountVectorizer()
countvect = cv.fit_transform(finaldata1['Text'].values)
print(type(countvect))
print(countvect.get_shape())
```

```
<class 'scipy.sparse.csr.csr_matrix'>
(100000, 61819)
```

In [11]:

```
score = finaldata1['Score']
print(score.shape)
```

```
(100000,)
```

In [12]:

```
score.head(5)
```

Out [12]:

```
425529    positive
494547    positive
147114    positive
441781    positive
276015    positive
Name: Score, dtype: object
```

RANDOM BASED SPLITTING OF DATA:

-> SIMPLE CROSS-VALIDATION

In [13]:

```
x_1,x_test,y_1,y_test = cross_validation.train_test_split(countvect,score,t
est_size=0.3,random_state=9)
```

Shape of Test and Train dataset:

-> The whole data is divided into training data contains 70% percent of data

-> 30 % of data for testing

In [14]:

```
print(x_1.get_shape())
print(x_test.get_shape())
print(y_1.count())
print(y_test.count())
```

```
(70000, 61819)
(30000, 61819)
70000
30000
```

Splitting the test data into test and cross-validation data:

-> Again the training data is divided into 70 and 30 ratio

-> This 70 % to train and 30 % for cross validation

In [15]:

```
x_train,x_cv,y_train,y_cv = cross_validation.train_test_split(x_1,y_1,test_
size=0.3,random_state=8)
```

In [16]:

```
print(x_train.get_shape())
print(x_cv.get_shape())
print(y_train.count())
print(y_cv.count())
```



```
(49000, 61819)
(21000, 61819)
49000
21000
```

In [18]:

```
for i in range(1,30,2):
    clas = KNeighborsClassifier(n_neighbors=i)
    clas.fit(x_train,y_train)
    pred = clas.predict(x_cv)
    acc = accuracy_score(y_cv,pred,normalize=True)*float(100)
    print('the accurate percentage of k = {} is {}'.format(i,acc))
```

```
the accurate percentage of k = 1 is 81.66190476190476
the accurate percentage of k = 3 is 84.04761904761905
the accurate percentage of k = 5 is 84.64285714285714
the accurate percentage of k = 7 is 84.63333333333334
the accurate percentage of k = 9 is 84.65238095238095
the accurate percentage of k = 11 is 84.67619047619047
the accurate percentage of k = 13 is 84.63809523809523
the accurate percentage of k = 15 is 84.57619047619048
the accurate percentage of k = 17 is 84.54285714285714
the accurate percentage of k = 19 is 84.4952380952381
the accurate percentage of k = 21 is 84.47619047619047
the accurate percentage of k = 23 is 84.46190476190476
the accurate percentage of k = 25 is 84.42380952380952
the accurate percentage of k = 27 is 84.39047619047619
the accurate percentage of k = 29 is 84.39999999999999
```

Note:

-> The highest accuracy 84.67 with cross-validation is occurred when k = 11

-> Using the same k-value and testing the accuracy on the test data set

In [22]:

```
clas = KNeighborsClassifier(n_neighbors = 11)
clas.fit(x_train,y_train)
pred = clas.predict(x_test)
acc = accuracy_score(y_test,pred,normalize=True)*float(100)
print("The accuracy on the test data set when k = {} is {}".format(11,acc))
```

The accuracy on the test data set when k = 11 is 84.75

Observation:

By using simple cross validation the highest accuracy achieved is 84.75 when k = 11

RANDOM BASED SPLITTING OF DATA:

K-FOLD CROSS VALIDATION

In [17]:

```
l = list(np.arange(1,30,2))
print(l)
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
```

In [24]:

```
cross_validation_scores = []
for k in l:
    classif = KNeighborsClassifier(n_neighbors=k)
    score =
cross_val_score(classif,x_train,y_train,cv=10,scoring='accuracy')
    cross_validation_scores.append(score.mean())
```

In [26]:

```
error = [1 - x for x in cross_validation_scores]
print("Cross Validation Accuracy Scores")
print(cross_validation_scores)
print("errors")
print(error)
```

Cross Validation Accuracy Scores

```
[0.8195103131824372, 0.841714120854399, 0.8463265325417131,
0.847489743752176, 0.8471020594378201, 0.8468366491810349,
0.8464489065576387, 0.8462652372455326, 0.8462040502375696,
0.8462244792263423, 0.8459795895868216, 0.8457959452643042,
0.8454489689899611, 0.8453673529932259, 0.8452857286666277]
```

errors


```
[0.18048968681756283, 0.158285879145601, 0.15367346745828692,
0.152510256247824, 0.15289794056217987, 0.15316335081896515,
0.15355109344236129, 0.1537347627544674, 0.15379594976243038,
0.15377552077365775, 0.15402041041317838, 0.15420405473569576,
0.15455103101003886, 0.15463264700677415, 0.15471427133337234]
```

In [27]:

```
mp.plot(l,cross_validation_scores)
mp.xlabel('k-values')
mp.ylabel('Cross Validation Accuracy')
mp.title("k vs CV Accuracy")
```

Out[27]:

Text(0.5,1,'k vs CV Accuracy')




k-values

In [28]:

```
mp.plot(1,error)
mp.xlabel('k-values')
mp.ylabel('error')
mp.title("k vs error")
```

Out[28]:

Text(0.5,1,'k vs error')



In [29]:

```
best_k = l[error.index(min(error))]  
print("the best value of k is {}".format(best_k))
```

the best value of k is 7

Observation:

-> By using k-fold cross validation we got best k value

-> From the plots we can say that error is less at k = 7 when compared with all of them

-> By using k = 7 designing the classifier with the training data

In [30]:

```
kclas = KNeighborsClassifier(n_neighbors=best_k)  
kclas.fit(x_train,y_train)  
predi = kclas.predict(x_test)  
accur = accuracy_score(y_test,predi,normalize=True)*float(100)  
print(accur)
```

84.83333333333334

Observation:

> Made the best k value to design the model with training data

-> Using the best k value to design the model with train data

-> The accuracy with the test data is 84.83

BAG OF WORDS:

TIME BASED SPLITTING OF DATA:

-> K-FOLD CROSS-VALIDATION:

-> TO FIND THE BEST K-VALUE WITH ACCURACY AS MEASURE

In order to perform this we have sort the data based on time

In [18]:

```
sorted_count_vect = finaldata1.sort_values("Time",axis=0,ascending=True,kind='quicksort',na_position='last',inplace=False)
```

Information about the sorted data:

-> Shape of the data

-> Dimensionality of the data

-> Attributes of the data

-> Sample of modified data

In [19]:

```
print(sorted_count_vect.shape)
print(sorted_count_vect.ndim)
print(sorted_count_vect.columns)
print(sorted_count_vect.head(5))
```

(100000, 10)

2


```
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
      'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')
```

	Id	ProductId	UserId	ProfileName	\
417839	451856	B00004CXX9	AIUWLEQ1ADEC5	Elizabeth Medina	
346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	
346094	374400	B00004CI84	A2DEE7F9XKP3ZR	jerome	
138017	149789	B00004S1C6	A1KXONFPU2XQ5K	Stephanie Manley	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	'
417839	0		0 positive	944092800	
346055	1		2 positive	944438400	
70688	0		0 positive	948672000	
346094	0		3 positive	959990400	
138017	26		28 positive	965779200	

	Summary \
417839	Entertainingl Funny!
346055	A modern day fairy tale
70688	A sure death for flies
346094	Research - Beatlejuice video - French version
138017	A must have!

	Text
417839	Beetlejuice is a well written movie ever...
346055	A twist of rumplestiskin captured on film, sta...
70688	I bought a few of these after my apartment was...
346094	I'm getting crazy.I'm looking for Beatlejuice ...
138017	These are easy to use, they do not make a mess...



In [20]:

```
count_vector = CountVectorizer()
```

In [21]:

```
cvr = count_vector.fit_transform(sorted_count_vect['Text'].values)
```

In [22]:

```
score = sorted_count_vect["Score"]
```

In [23]:

```
print(cvr.shape)
print(score.shape)
```

```
(100000, 61819)
(100000,)
```

In [24]:

```
x_train = cvr[0:70000]
x_test = cvr[70000:100000]
y_train = score[0:70000]
y_test = score[70000:100000]
```

Splitting the sampled data into test and train:

-> Training data contains 70 percent of data

-> Test data contains 30 percent of data

In [25]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(70000, 61819)
(30000, 61819)
(70000,)
(30000,)
```

TIME BASED SPLITTING :

K-FOLD CROSS VALIDATION

In [26]:

```
l = list(np.arange(1,30,2))
print(l)
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
```

In [54]:

```
cross_validation_scores = []
for k in l:
    classif = KNeighborsClassifier(n_neighbors=k)
    score =
cross_val_score(classif,x_train,y_train,cv=10,scoring='accuracy')
    cross_validation_scores.append(score.mean())
```

In [55]:

```
error = [1 - x for x in cross_validation_scores]
print("Cross Validation Accuracy Scores")
print(cross_validation_scores)
print("errors")
print(error)
```

Cross Validation Accuracy Scores

```
[0.8229853780201098, 0.8476284336250701, 0.8520285439189499,
0.853171507207582, 0.8533857908854243, 0.8530286500618093,
0.8529143806717221, 0.8525715153586024, 0.8523857500486887,
0.8521714377994171, 0.8521000235122453, 0.8514714561524788,
0.8513571500276969, 0.8512428520661807, 0.8513285663536442]
```

errors


```
[0.17701462197989015, 0.15237156637492988, 0.1479714560810501,
0.14682849279241805, 0.14661420911457568, 0.14697134993819072,
0.1470856193282779, 0.14742848464139757, 0.14761424995131134,
0.14782856220058294, 0.14789997648775466, 0.14852854384752123,
0.14864284997230315, 0.14875714793381933, 0.14867143364635582]
```


In [56]:

```
mp.plot(l,cross_validation_scores)
mp.xlabel('k-values')
mp.ylabel('Cross Validation Accuracy')
mp.title("k vs CV Accuracy")
```

Out[56]:

Text(0.5,1,'k vs CV Accuracy')






In [57]:

```
mp.plot(1,error)
mp.xlabel('k-values')
mp.ylabel('error')
mp.title("k vs error")
```

Out[57]:

Text(0.5,1,'k vs error')



In [58]:

```
best_k = 1[error.index(min(error))]
print("the best value of k is {}".format(best_k))
```

the best value of k is 9

Observation:

- > By using k-fold cross validation we got best k value as 9
- > From the plots we can say that error is less at k = 9 when compared with all of them
- > By using k = 9 designing the classifier with the training data

In [59]:

```
kclas = KNeighborsClassifier(n_neighbors=best_k)
kclas.fit(x_train,v_train)
```

```

predi = kclas.predict(x_test)
accur = accuracy_score(y_test,predi,normalize=True)*float(100)
print(accur)

```

83.22666666666667

Observation:

-> Using the best k value to design the model with train data

-> The accuracy with the test data is 83.2266

TFIDF:

TIME BASED SPLITTING OF DATA:

-> K-FOLD CROSS-VALIDATION:

-> TO FIND THE BEST K-VALUE WITH ACCURACY AS MEASURE

In [27]:

```

tfidf_data = finaldata1.sort_values("Time",axis=0,ascending=True,kind='quic
ksort',na_position='last',inplace=False)

```

Information about the sorted data:

-> Shape of the data

-> Dimensionality of the data

-> Attributes if the data

-> Sample of modified data

In [28]:

```

print(tfidf_data.shape)
print(tfidf_data.ndim)
print(tfidf_data.columns)
print(tfidf_data.head(5))

```

(100000, 10)

2

```

Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
       'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')

```

	Id	ProductId	UserId	ProfileName	\
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	
346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	
346094	374400	B00004CI84	A2DEE7F9XKP3ZR	jerome	
138017	149789	B00004S1C6	A1KXONFPU2XQ5K	Stephanie Manley	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time
417839	0	0	positive	944092800
346055	1	2	positive	944438400
70688	0	0	positive	948672000
346094	0	3	positive	959990400
138017	26	28	positive	965779200

	Summary
417839	Entertainingl Funny!
346055	A modern day fairy tale
70688	A sure death for flies
346094	Research - Beatlejuice video - French version
138017	A must have!

	Text
417839	Beetlejuice is a well written movie ever...
346055	A twist of rumplestiskin captured on film, sta...
70688	I bought a few of these after my apartment was...
346094	I'm getting crazy.I'm looking for Beatlejuice ...
138017	These are easy to use, they do not make a mess...

In [29]:

```
tfid_vect = TfidfVectorizer(ngram_range=(1,2))
```

In [30]:

```
tfid_vect_data = tfidf_vect.fit_transform(tfidf_data["Text"].values)
```

In [31]:

```
print(tfid_vect_data.shape)
print(tfidf_data["Score"].shape)
```

```
(100000, 1287433)
(100000,)
```

Splitting the sampled data into test and train:

-> Training data contains 70 percent of data

-> Test data contains 30 percent of data

In [32]:

```
x_train = tfidf_vect_data[0:70000]
x_test = tfidf_vect_data[70000:100000]
y_train = tfidf_data['Score'][0:70000]
y_test = tfidf_data['Score'][70000:100000]
```

In [33]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(70000, 1287433)
(30000, 1287433)
```

```
(70000,)
(30000,)
```

TIME BASED SPLITTING:

K-FOLD CROSS VALIDATION

In [34]:

```
l = list(np.arange(1,30,2))
print(l)
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
```

In [71]:

```
cross_validation_scores = []
for k in l:
    clf = KNeighborsClassifier(n_neighbors=k)
    score =
cross_val_score(clf,x_train,y_train,cv=10,scoring='accuracy')
    cross_validation_scores.append(score.mean())
```

In [72]:

```
error = [1 - x for x in cross_validation_scores]
print("Cross Validation Accuracy Scores")
print(cross_validation_scores)
print("errors")
print(error)
```


```
Cross Validation Accuracy Scores
[0.8475139805615098, 0.8578856093446042, 0.863228519657725,
0.8653428870478141, 0.8665142789084547, 0.867514401377845,
0.8669716176932983, 0.866400127885717, 0.8660715829810528,
0.8658715625688075, 0.8651001707163302, 0.864414480906418,
0.8641001584510237, 0.864042946204958, 0.8634000992530633]
errors
[0.15248601943849016, 0.14211439065539577, 0.13677148034227504,
0.13465711295218585, 0.13348572109154533, 0.13248559862215503,
0.13302838230670166, 0.13359987211428304, 0.1339284170189472,
0.13412843743119252, 0.13489982928366984, 0.13558551909358196,
0.13589984154897627, 0.13595705379504197, 0.13659990074693673]
```

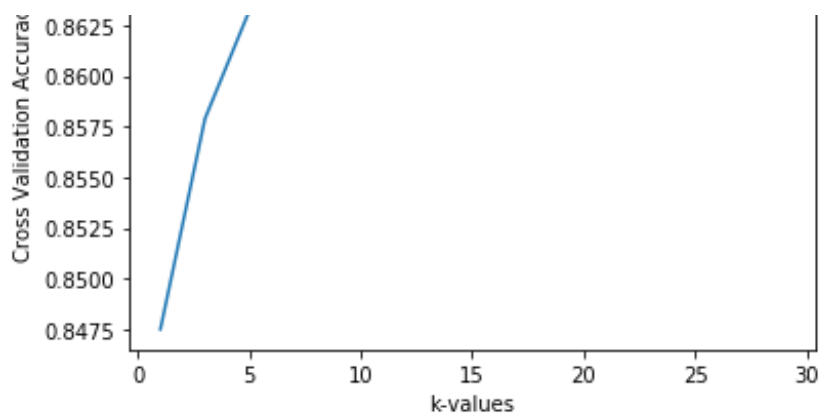
In [73]:

```
mp.plot(l,cross_validation_scores)
mp.xlabel('k-values')
mp.ylabel('Cross Validation Accuracy')
mp.title("k vs CV Accuracy")
```

Out[73]:

```
Text(0.5,1,'k vs CV Accuracy')
```






In [77]:

```
mp.plot(1,error)
mp.xlabel('k-values')
mp.ylabel('error')
mp.title("k vs Error")
```

Out[77]:

Text(0.5,1,'k vs Error')



In [78]:

```
best_k = 1[error.index(min(error)) ]
print("the best value of k is {}".format(best_k))
```

the best value of k is 11

Observation:

-> By using k-fold cross validation we got best k value as 11

-> From the plots we can say that error is less at k = 11 when compared with all of them

-> By using k = 11 designing the classifier with the training data

In [79]:

```
kclas = KNeighborsClassifier(n_neighbors=best_k)
kclas.fit(x_train,y_train)
predi = kclas.predict(x_test)
accur = accuracy_score(y_test,predi,normalize=True)*float(100)
print(accur)
```

84.98333333333333

Observation:

-> Using the best k value to design the model with train data

-> The accuracy with the test data is 84.9833

CONSTRUCTING VECTOR REPRESENTATION OF EACH IN THE DATA BY USING WORD2VEC

In [35]:

```
import gensim
from gensim.models import word2vec
```

-> Importing the required libraries

-> Functions to clean the sentences

-> Constructing the word2vec from the sample subset data

In [36]:

```
import re
def cleanhtml(sentence):
    clean = re.compile("<.*?>")
    cleantext = re.sub(clean, " ",sentence)
    return cleantext
def cleanpunct(sentence):
    cleanr = re.sub(r"[?|!|\\|'|#|.|.|,|)|(|/|)",r' ',sentence)
    return cleanr
```

In [37]:

```
sorted_w2vec = finaldata1.sort_values("Time",axis=0,ascending=True,kind='quicksort',na_position='last',inplace=False)
```

Information about the sorted data:

-> Shape of the data

-> Dimensionality of the data

-> Attributes of the data

-> Sample of modified data

In [38]:

```
print(sorted_w2vec.shape)
print(sorted_w2vec.ndim)
print(sorted_w2vec.columns)
print(sorted_w2vec.head(5))
```

```
(100000, 10)
```

```
2
```

```
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
      'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],
      dtype='object')
```

	Id	ProductId	UserId	ProfileName	\
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	
346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	
346094	374400	B00004CI84	A2DEE7F9XKP3ZR	jerome	
138017	149789	B00004S1C6	A1KXONFPU2XQ5K	Stephanie Manley	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	'
417839	0	0	positive	944092800	
346055	1	2	positive	944438400	
70688	0	0	positive	948672000	
346094	0	3	positive	959990400	
138017	26	28	positive	965779200	

	Summary	\
417839	Entertainingl Funny!	
346055	A modern day fairy tale	
70688	A sure death for flies	
346094	Research - Beatlejuice video - French version	
138017	A must have!	

	Text
417839	Beetlejuice is a well written movie ever...
346055	A twist of rumplestiskin captured on film, sta...
70688	I bought a few of these after my apartment was...
346094	I'm getting crazy.I'm looking for Beatlejuice ...
138017	These are easy to use, they do not make a mess...

In [39]:

```
i=0
sentences_list=[]
for sent in sorted_w2vec['Text'].values:
    filtered_sentences = []
    sent = cleanhtml(sent)
    for w in sent.split():
        for cleanedwords in cleanpunct(w).split():
            if(cleanedwords.isalpha()):
                filtered_sentences.append(cleanedwords.lower())
    sentences_list.append(filtered_sentences)
```

In [40]:

```
print(len(sentences_list))
print(type(sentences_list))
```

```
100000
```

```
<class 'list'>
```

```
In [41]:
```

```
w2vmodel =  
gensim.models.Word2Vec(sentences_list,min_count=4,size=40,workers=4)
```

-> Most similar word

-> Similarity between the words

-> Dimensionality representation of a word

```
In [42]:
```

```
print(w2vmodel.most_similar("hello"))  
print(w2vmodel.similarity("hello",'hi'))  
print(w2vmodel.wv['hello'])
```

```
[('mickey', 0.7368489503860474), ('anyhow', 0.7226568460464478),  
( 'focaccia', 0.709900975227356), ('fern', 0.7098590731620789), ('mike', 0.6  
99766993522644), ('boomer', 0.6946572661399841), ('themed',  
0.6860719919204712), ('merry', 0.6846867799758911), ('claxton', 0.683440923  
6907959), ('samoas', 0.6817569732666016)]  
0.5446572411054558  
[ 0.29999846 -0.19882171 -0.41949722 -0.4891025 -0.3457285  0.03067495  
  0.33179036  0.11951371 -0.3910432 -0.33136272 -0.08102001  0.09156115  
  0.2493193  0.2969718  0.23660749 -0.362925 -0.38329345  0.02786745  
  0.2396555 -0.12413767  0.32540685 -0.04338589  0.27390665  0.233951  
 -0.24852224  0.18082564  0.3852502  0.25250012 -0.04117521  0.1665598  
 -0.1021607  0.06984844  0.10276769  0.21571645  0.4212314  0.19228005  
  0.01020537  0.06221204 -0.18181211  0.15039672]
```

```
/Users/vthumati/anaconda3/lib/python3.6/site-  
packages/ipykernel_launcher.py:1: DeprecationWarning: Call to deprecated `m  
ost_similar` (Method will be removed in 4.0.0, use self.wv.most_similar() i  
nstead).
```

```
"""Entry point for launching an IPython kernel.
```

```
/Users/vthumati/anaconda3/lib/python3.6/site-  
packages/ipykernel_launcher.py:2: DeprecationWarning: Call to deprecated `s  
imilarity` (Method will be removed in 4.0.0, use self.wv.similarity() inste  
ad).
```

Observation:

-> We have constructed the vector representation of each word

-> Using this model to construct vector representation of each sentence in average word2vec and tfidf-word2vec

AVERAGE WORD2VEC

-> Here i am using the word2vec model to construct vector representation of each sentence

-> Sorting the data based on the time

In [43]:

```
sent_vectors = []
for sent in sentences_list:
    sent_vec = np.zeros(40)
    cnt=0
    for word in sent:
        try:
            vec = w2vmodel.wv[word]
            sent_vec += vec
            cnt += 1
        except:
            pass
    sent_vec /= cnt
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[88888]))
```

100000
40

In [44]:

```
np.isnan(sent_vectors).any()
```

Out[44]:

False

In [45]:

```
type(sent_vectors)
```

Out[45]:

list

In [46]:

```
sent_vectors = np.nan_to_num(sent_vectors)
```

In [47]:

```
type(sent_vectors)
```

Out[47]:

numpy.ndarray

In [48]:

```
sent_vectors.shape
```

Out[48]:

(100000, 40)

In [49]:

```
x_train = sent_vectors[0:70000]
x_test = sent_vectors[70000:100000]
y_train = sorted w2vec['Score'][0:70000]
```

```
y_test = sorted_w2vec['Score'][70000:100000]
```

In [50]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(70000, 40)
(30000, 40)
(70000,)
(30000,)
```

In [51]:

```
print(l)
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
```

In [52]:

```
from sklearn.cross_validation import cross_val_score
```

In [53]:

```
cross_validation_scores = []
for k in l:
    classif = KNeighborsClassifier(n_neighbors=k)
    score =
cross_val_score(classif,x_train,y_train,cv=10,scoring='accuracy')
    cross_validation_scores.append(score.mean())
```

In [54]:

```
error = [1 - x for x in cross_validation_scores]
print("Cross Validation Accuracy Scores")
print(cross_validation_scores)
print("errors")
print(error)
```

Cross Validation Accuracy Scores


```
[0.8293284544422717, 0.8614428912836741, 0.8692571606994173,
0.8718571300880464, 0.873614285205831, 0.873585701530612,
0.8744857484731785, 0.8742428852075808, 0.8747000239784262,
0.8746142688775507, 0.8741142811218658, 0.8743285770411079,
0.8747142913268224, 0.8742857015282797, 0.8737857035705536]
errors
[0.17067154555772834, 0.13855710871632587, 0.13074283930058272,
0.12814286991195356, 0.12638571479416905, 0.126414298469388,
0.1255142515268215, 0.12575711479241924, 0.12529997602157383,
0.12538573112244933, 0.12588571887813416, 0.1256714229588921,
0.12528570867317756, 0.12571429847172033, 0.12621429642944637]
```

In [55]:

```
mp.plot(l,cross_validation_scores)
mp.xlabel('k-values')
mp.ylabel('Cross Validation Accuracy')
mp.title("k vs CV Accuracy")
```


Out[55]:

Text(0.5,1,'k vs CV Accuracy')




In [56]:

```
mp.plot(1,error)
mp.xlabel('k-values')
mp.ylabel('error')
mp.title("k vs Error")
```

Out[56]:

Text(0.5,1,'k vs Error')



In [57]:

```
best_k = 1[error.index(min(error))]  
print("the best value of k is {}".format(best_k))
```

the best value of k is 25

Observation:

-> By using k-fold cross validation we got best k value as 25

-> From the plots we can say that error is less at $k = 25$ when compared with all of them

-> By using $k = 25$ designing the classifier with the training data

In [58]:

```
kclas = KNeighborsClassifier(n_neighbors=best_k)
kclas.fit(x_train,y_train)
predi = kclas.predict(x_test)
accur = accuracy_score(y_test,predi,normalize=True)*float(100)
print(accur)
```

85.71666666666667

Observation:

-> Using the best k value to design the model with train data

-> The accuracy with the test data is 85.71

TFIDF-WORD2VEC:

Here i am using the TFIDF-WORD2VEC with the help of word2vec model to construct vector representation of each sentence

In [53]:

```
tfidf_feat = tfidf_vect.get_feature_names()

tfidf_sent_vectors = []

row=0
for sent in sentences_list:
    sent_vec = np.zeros(40)
    sum =0
    for word in sent:
        try:
            vec = w2v_model.wv[word]

            tfidf = tfidf_vect_data[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            sum += tfidf
        except:
            pass
    sent_vec /= sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
```

/Users/vthumati/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:20: RuntimeWarning: invalid value encountered in true_divide

In [54]:

```
print(len(tfidf_sent_vectors))  
print(len(tfidf_sent_vectors[99999]))
```

```
100000  
40
```

In [55]:

```
type(tfidf_sent_vectors)
```

Out[55]:

```
list
```

In [56]:

```
np.isnan(tfidf_sent_vectors).any()
```

Out[56]:

```
True
```

In [57]:

```
tfidf_sent_vectors = np.nan_to_num(tfidf_sent_vectors)
```

In [58]:

```
np.isnan(tfidf_sent_vectors).any()
```

Out[58]:

```
False
```

In [60]:

```
x_train = tfidf_sent_vectors[0:70000]  
x_test = tfidf_sent_vectors[70000:]  
y_train = sorted_w2vec['Score'][0:70000]  
y_test = sorted_w2vec['Score'][70000:]
```

In [62]:

```
print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

```
(70000, 40)  
(30000, 40)  
(70000,)  
(30000,)
```

In [63]:

```
print(1)
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]
```

In [64]:

```
from sklearn.cross_validation import cross_val_score
```

In [65]:

```
cross_validation_scores = []
for k in l:
    clf = KNeighborsClassifier(n_neighbors=k)
    score =
cross_val_score(clf,x_train,y_train,cv=10,scoring='accuracy')
    cross_validation_scores.append(score.mean())
```

In [66]:

```
error = [1 - x for x in cross_validation_scores]
print("Cross Validation Accuracy Scores")
print(cross_validation_scores)
print("errors")
print(error)
```

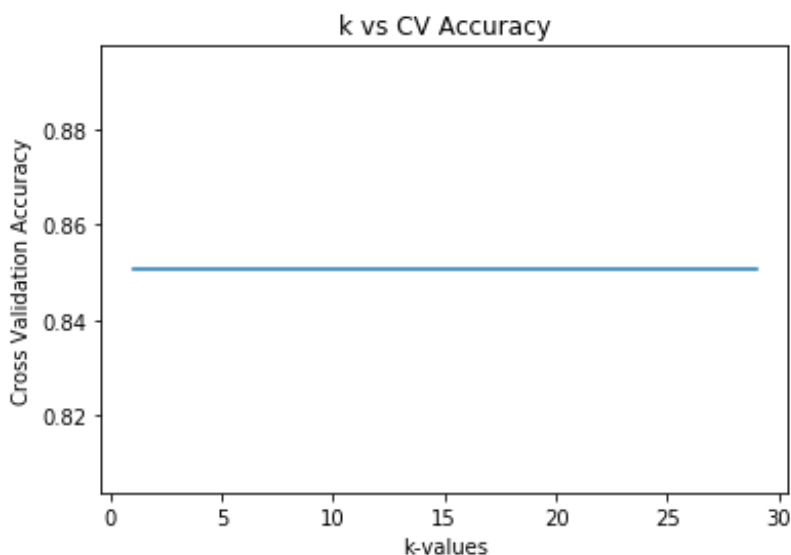
```
Cross Validation Accuracy Scores
[0.8506285771533528, 0.8506285771533528, 0.8506285771533528,
0.8506285771533528, 0.8506285771533528, 0.8506285771533528,
0.8506285771533528, 0.8506285771533528, 0.8506285771533528,
0.8506285771533528, 0.8506285771533528, 0.8506285771533528,
0.8506285771533528, 0.8506285771533528, 0.8506285771533528]
errors
[0.14937142284664717, 0.14937142284664717, 0.14937142284664717,
0.14937142284664717, 0.14937142284664717, 0.14937142284664717,
0.14937142284664717, 0.14937142284664717, 0.14937142284664717,
0.14937142284664717, 0.14937142284664717, 0.14937142284664717,
0.14937142284664717, 0.14937142284664717, 0.14937142284664717]
```

In [67]:

```
mp.plot(l,cross_validation_scores)
mp.xlabel('k-values')
mp.ylabel('Cross Validation Accuracy')
mp.title("k vs CV Accuracy")
```

Out[67]:

Text(0.5,1,'k vs CV Accuracy')




In [68]:

```
mp.plot(1,error)
mp.xlabel('k-values')
mp.ylabel('error')
mp.title("k vs Error")
```

Out[68]:

Text(0.5,1,'k vs Error')



In [69]:

```
best_k = 1[error.index(min(error))]  
print("the best value of k is {}".format(best_k))
```

the best value of k is 1

In [70]:

```
kclas = KNeighborsClassifier(n_neighbors=best_k)  
kclas.fit(x_train,y_train)  
predi = kclas.predict(x_test)  
accur = accuracy_score(y_test,predi,normalize=True)*float(100)  
print(accur)
```

82.47

Observation:

-> Using the best k value which is 1 to design the model with train data

-> The accuracy with the test data is 82.47

CONCLUSION:

RANDOM BASED SPLITTING:

-> THIS IS PERFORMED ONLY ON BAG OF WORDS TECHNIQUE

-> BOTH SIMPLE CROSS VALIDATION AND 10-FOLD CROSS VALIDATION ARE

PERFORMED ON THIS
TECHNIQUE

-> IN SIMPLE CROSS VALIDATION THE BEST K - VALUE IS 11 AND THE ACCURACY IS 84.67

-> IN 10-FOLD CROSS VALIDATION THE BEST K - VALUE IS 7 AND THE ACCURACY IS 84.83

TIME BASED SPLITTING:

-> THIS IS PERFORMED ON BAG OF WORDS, TFIDF, AVERAGE WORD2VEC

USING 10 - FOLD CROSS VALIDATION:

BAG OF WORDS:

-> THE BEST K - VALUE IS 9 AND THE ACCURACY IS 83.22

TFIDF:

-> THE BEST K - VALUE IS 11 AND THE ACCURACY IS 84.98

AVERAGE WORD2VEC:

-> THE BEST K - VALUE IS 25 AND THE ACCURACY IS 85.71

TFIDF-WORD2VEC:

-> THE BEST K - VALUE IS 1 AND ACCURACY IS 82.47