

Information about data:

- > The MNIST dataset is in built in Tensor Flow Tutorials or we can get from Kaggle
- > The dataset contains of Hand written digits and their corresponding class lables

Objective:

- > To perform Multi Layer Perceptron on the MNIST dataset
- > To introduce different optimization algorithms like Gradient Descent, Stochastic Gradient Descent, Adadelata, Adam
- > To introduce different initializations of weights like Xavier, He
- > To try different activation functions like RELU, Sigmoid, Tanh
- > To introduce Batch Normalization
- > To introduce Drop out
- > To analyze the weights for each model

-> Importing the required libraries

-> Importing the in built dataset from Tensor Flow

In [3]:

```
import tensorflow as tf
import matplotlib.pyplot as mp
import numpy as np
from tensorflow.examples.tutorials.mnist import input_data
data = input_data.read_data_sets("MNIST_data/", one_hot = True)
```

WARNING:tensorflow:From <ipython-input-3-6c9780a2e235>:5: read\_data\_sets (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use alternatives such as official/mnist/dataset.py from tensorflow/models.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:260: maybe\_download (from tensorflow.contrib.learn.python.learn.datasets.base) is deprecated and will be removed in a future version.

Instructions for updating:

Please write your own downloading logic.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/base.py:252: \_internal\_retry.<locals>.wrap.<locals>.wrapped\_fn (from tensorflow.contrib.learn.python.learn.datasets.base) is deprecated and will be removed in a future version.

Instructions for updating:

```

Please use urllib or similar directly.
Successfully downloaded train-images-idx3-ubyte.gz 9912422 bytes.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:262: extra
ct_images (from tensorflow.contrib.learn.python.learn.datasets.mnist) is de
precated and will be removed in a future version.
Instructions for updating:
Please use tf.data to implement this functionality.
Extracting MNIST_data/train-images-idx3-ubyte.gz
Successfully downloaded train-labels-idx1-ubyte.gz 28881 bytes.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:267: extra
ct_labels (from tensorflow.contrib.learn.python.learn.datasets.mnist) is de
precated and will be removed in a future version.
Instructions for updating:
Please use tf.data to implement this functionality.
Extracting MNIST_data/train-labels-idx1-ubyte.gz
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:110: dense
_to_one_hot (from tensorflow.contrib.learn.python.learn.datasets.mnist) is
deprecated and will be removed in a future version.
Instructions for updating:
Please use tf.one_hot on tensors.
Successfully downloaded t10k-images-idx3-ubyte.gz 1648877 bytes.
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Successfully downloaded t10k-labels-idx1-ubyte.gz 4542 bytes.
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:290: DataS
et.__init__ (from tensorflow.contrib.learn.python.learn.datasets.mnist) is
deprecated and will be removed in a future version.
Instructions for updating:
Please use alternatives such as official/mnist/dataset.py from
tensorflow/models.

```

-> Dynamic plot function

In [0]:

```

def dynamic_plot(x, y, y_1, ax, ticks, title, colors=['b']):
    ax.plot(x, y, 'b', label="Train Loss")
    ax.plot(x, y_1, 'r', label="Test Loss")
    if len(x)==1:
        mp.legend()
        mp.title(title)
    mp.yticks(ticks)
    fig.canvas.draw()

```

Train and Test data

In [5]:

```

print(data.train.images.shape)
print(data.test.images.shape)
print(data.train.labels.shape)
print(data.test.labels.shape)

```

```
print(data.train.images.shape,
```

```
(55000, 784)  
(10000, 784)  
(55000, 10)  
(10000, 10)
```

In [6]:

```
print('The number of data point in training are:',data.train.images.shape[0]  
,"and number of features for each data point are",data.train.images.shape[1])  
print('The number of data point to test are:',data.test.images.shape[0],"a  
nd number of features for each data point are",data.test.images.shape[1])  
print("The shape of the class labels in the train and test dataset after ap  
pling One Hot Encoding on the class labels")  
print("The number of class labels in training are:",data.train.labels.shape  
[0],"and each class label is of:",data.train.labels.shape[1],"dimensions")  
print("The number of class labels in training are:",data.test.labels.shape[  
0],"and each class label is of:",data.test.labels.shape[1],"dimensions")
```

The number of data point in training are: 55000 and number of features for each data point are 784

The number of data point to test are: 10000 and number of features for each data point are 784

The shape of the class labels in the train and test dataset after applying One Hot Encoding on the class labels

The number of class labels in training are: 55000 and each class label is of: 10 dimensions

The number of class labels in training are: 10000 and each class label is of: 10 dimensions

#### NETWORK PARAMETERS:

-> Number of hidden layers

-> Number of Neurons in each Layer

-> The input is of 784 Dimension

-> The output is a SOFTMAX classifier

with 10 class labels

In [0]:

```
n_hidden_1 = 512  
n_hidden_2 = 128  
n_input = 784  
n_output = 10
```

#### Variables and Place holders

In [0]:

```
x = tf.placeholder(tf.float32, [None, 784])
```

```

y_ = tf.placeholder(tf.float32, [None,10])
prob = tf.placeholder(tf.float32) ##### place holder for drop out of input
prob_input = tf.placeholder(tf.float32) ##### place holder for drop out of
output
keep_prob = tf.placeholder(tf.float32)
keep_prob_input = tf.placeholder(tf.float32)

```

Initializing Weights and Bias using different initialization algorithms like Xavier/Glorot, He

-> Xavier/Glorot weight initialization for SGD

-> He weight initialization for ReLU

In [0]:

```

sgd_weights = {
    'h1' : tf.Variable(tf.random_normal([n_input,n_hidden_1],stddev=0.039,mean=0)), # 784*512 # sqrt(2/(784+512)) = 0.039
    'h2' : tf.Variable(tf.random_normal([n_hidden_1,n_hidden_2],stddev=0.055,mean=0)), # 512*128 # sqrt(2/(512+128)) = 0.055
    'out': tf.Variable(tf.random_normal([n_hidden_2,n_output],stddev=0.120,mean = 0)) # 128*10 #sqrt(2/(128+10)) = 0.120
}

relu_weights = {
    'h1' : tf.Variable(tf.random_normal([n_input,n_hidden_1],stddev=0.062,mean=0)), # 784*512 # sqrt(2/(784+1)) = 0.062
    'h2' : tf.Variable(tf.random_normal([n_hidden_1,n_hidden_2],stddev=0.125,mean=0)), # 512*128 # sqrt(2/(512+1)) = 0.125
    'out': tf.Variable(tf.random_normal([n_hidden_2,n_output],stddev=0.120,mean = 0)) # 128*10 #sqrt(2/(128+1)) = 0.120
}

biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_output]))
}

```

## PARAMETERS

In [0]:

```

training_epochs = 25
learning_rate = 0.05
batch_size = 1000
display_step = 1

```

Model 1: INPUT (784) - SIGMOID(512) - SIGMOID(128) - SIGMOID(output 10)

-> Sigmoid Activation for hidden layers

In [0]:

```
def mulit_layer_perceptron(x,weights,biases):
    print('Input layer')
    print("X:", x.get_shape(), "W[h1]:", weights['h1'].get_shape(), 'b[h1]:',
biases['b1'].get_shape())

    print('Hidden layer 1 with Sigmoid Activation')
    layer_1 = tf.add(tf.matmul(x,weights['h1']),biases['b1'])
    layer_1 = tf.nn.sigmoid(layer_1)
    print("Layer_1:", layer_1.get_shape(), 'W[h2]:', weights['h2'].get_shape,
'b[b1]:', biases['b2'].get_shape())

    print("Hidden layer 2 with Sigmoid Activation")
    layer_2 = tf.add(tf.matmul(layer_1,weights['h2']),biases['b2'])
    layer_2 = tf.nn.sigmoid(layer_2)
    print("Layer_2:", layer_2.get_shape(), 'W[out]:', weights['out'].get_shape,
'e, 'b3:', biases['out'].get_shape())

    print("Output layer")
    output_layer = tf.matmul(layer_2,weights['out']) + biases['out']
    output_layer = tf.nn.sigmoid(output_layer)
    print('Output Layer:', output_layer.get_shape())

    return output_layer
```

#### MODEL 1 + ADAM OPTIMIZER

In [33]:

```
y_sgd = mulit_layer_perceptron(x,sgd_weights,biases)
cost_sgd = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_
sgd, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_sgd)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_sgd)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([adam_optimizer,cost_sgd,sgd_weights],feed_dict = {x:
xs_batch, y_:ys_batch})
```

```

        train_avg_cost += c/number_of_batches
        c = sess.run(cost_sgd, feed_dict = {x: data.test.images, y_: data.test
.labels})
        test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid(512)-sigmoid(128)-sigmoid(output)-AdamOptimizer")

    if epoch%display_step == 0:
        print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_a
vg_cost), "test cost={:.9f}".format(test_avg_cost))
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid(512)-sigmoid(128)-sigmoid(output)-AdamOptimizer")
        correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labe
ls}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with Sigmoid Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_9:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with Sigmoid Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_10:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)

Output layer

Output Layer: (?, 10)

Epoch: 0001 train cost=2.380851776 test cost=2.383974218

Epoch: 0002 train cost=2.390954169 test cost=2.386933439

Epoch: 0003 train cost=2.421044302 test cost=2.421044510

Epoch: 0004 train cost=2.416726303 test cost=2.419485166

Epoch: 0005 train cost=2.420288840 test cost=2.419811010

Epoch: 0006 train cost=2.414657515 test cost=2.416858786

Epoch: 0007 train cost=2.418963445 test cost=2.416482449

Epoch: 0008 train cost=2.414868160 test cost=2.416482449

Epoch: 0009 train cost=2.417286340 test cost=2.416482449

Epoch: 0010 train cost=2.416559072 test cost=2.416482449

Epoch: 0011 train cost=2.416633970 test cost=2.416482449

Epoch: 0012 train cost=2.416597613 test cost=2.416482449

Epoch: 0013 train cost=2.416495098 test cost=2.416482449

Epoch: 0014 train cost=2.414902340 test cost=2.416482449

Epoch: 0015 train cost=2.418383806 test cost=2.416482449

Epoch: 0016 train cost=2.415936505 test cost=2.416482449

Epoch: 0017 train cost=2.415633960 test cost=2.416482449

Epoch: 0018 train cost=2.417508897 test cost=2.416482449

Epoch: 0019 train cost=2.416322699 test cost=2.416482449

Epoch: 0020 train cost=2.416581622 test cost=2.416482449

Epoch: 0021 train cost=2.415613604 test cost=2.416482449


Epoch: 0022 train cost=2.417015800 test cost=2.416482449

Epoch: 0023 train cost=2.415947450 test cost=2.416482449

Epoch: 0024 train cost=2.416775062 test cost=2.416482449

Epoch: 0025 train cost=2.416436165 test cost=2.416482449

Accuracy: 0.1032



#### DISTRIBUTION OF WEIGHTS AT THE END OF TRAINING

In [36]:

```
import seaborn as s


h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  violin_data = remove_na(group_data)
```



MODEL 1 + GRADIENT DESCENT OPTIMIZER

In [40]:

```
y_sgd = mulit_layer_perceptron(x,sgd_weights,biases)
cost_sgd = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_
sgd, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_sgd)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_sgd)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([sgd_optimizer,cost_sgd,sgd_weights],feed_dict = {x:x
s_batch, y_:ys_batch})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_sgd, feed_dict = {x: data.test.images, y_: data.test
.labels})
            test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid(512)-sigmoid(128)-sigmoid(output)-Gradient Descent O
ptimizer")
```



```

if epoch%display_step == 0:
    print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))
    dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.5, 2.5, step=0.1), "input-sigmoid(512)-sigmoid(128)-sigmoid(output)-Gradient Descent Optimizer")
    correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labels}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with Sigmoid Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_9:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with Sigmoid Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_10:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)

Output layer

Output Layer: (?, 10)

Epoch: 0001 train cost=2.043193551 test cost=2.029072599

Epoch: 0002 train cost=1.828018275 test cost=1.818836763

Epoch: 0003 train cost=1.758738297 test cost=1.750763631

Epoch: 0004 train cost=1.723038485 test cost=1.715592033

Epoch: 0005 train cost=1.699364166 test cost=1.692855772

Epoch: 0006 train cost=1.684273349 test cost=1.676451995

Epoch: 0007 train cost=1.670672937 test cost=1.663909405

Epoch: 0008 train cost=1.660386177 test cost=1.653847081

Epoch: 0009 train cost=1.652663500 test cost=1.645478160

Epoch: 0010 train cost=1.644776596 test cost=1.638397306

Epoch: 0011 train cost=1.640572270 test cost=1.632387510

Epoch: 0012 train cost=1.633569223 test cost=1.627141695

Epoch: 0013 train cost=1.629938219 test cost=1.622509373

Epoch: 0014 train cost=1.625114491 test cost=1.618449651

Epoch: 0015 train cost=1.621762742 test cost=1.614725340

Epoch: 0016 train cost=1.618119214 test cost=1.611504769

Epoch: 0017 train cost=1.615162013 test cost=1.608589630

Epoch: 0018 train cost=1.612940318 test cost=1.605904512

Epoch: 0019 train cost=1.610706869 test cost=1.603522582

Epoch: 0020 train cost=1.606590884 test cost=1.601215137

Epoch: 0021 train cost=1.607143675 test cost=1.599224574


Epoch: 0022 train cost=1.602653252 test cost=1.597340612


Epoch: 0023 train cost=1.601170429 test cost=1.595592232

Epoch: 0024 train cost=1.601869932 test cost=1.594005286

Epoch: 0025 train cost=1.598346431 test cost=1.592462624

Accuracy: 0.8936





In [41]:

```
import seaborn as s

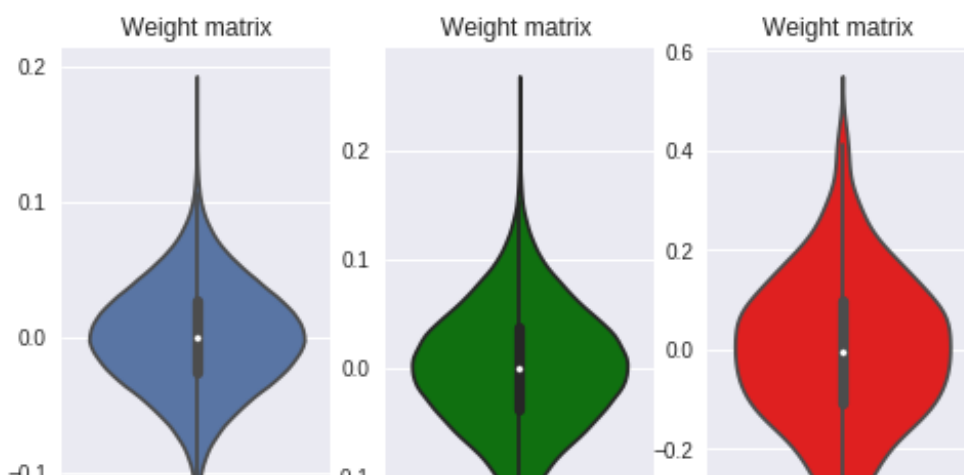
h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)


fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 kde\_data = remove\_na(group\_data)  
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 violin\_data = remove\_na(group\_data)





MODEL 2: INPUT (784) - ReLu(512) - ReLu(128) - ReLu(output  
10)

-> The activation functions are ReLu

In [0]:

```
def mulit_layer_perceptron(x,weights,biases):
    print('Input layer')
    print("X:", x.get_shape(), "W[h1]:", weights['h1'].get_shape(), 'b[h1]:',
    biases['b1'].get_shape())

    print('Hidden layer 1 with RELU Activation')
    layer_1= tf.add(tf.matmul(x,weights['h1']),biases['b1'])
    layer_1 = tf.nn.relu(layer_1)
    print("Layer_1:", layer_1.get_shape(), 'W[h2]:', weights['h2'].get_shape,
    'b[b1]:', biases['b2'].get_shape())

    print("Hidden layer 2 with RELU Activation")
    layer_2 = tf.add(tf.matmul(layer_1,weights['h2']),biases['b2'])
    layer_2 = tf.nn.relu(layer_2)
    print("Layer_2:", layer_2.get_shape(), 'W[out]:', weights['out'].get_shap
    e, 'b3:', biases['out'].get_shape())

    print("Output layer")
    output_layer = tf.matmul(layer_2,weights['out']) + biases['out']
    output_layer = tf.nn.relu(output_layer)
    print('Output Layer:', output_layer.get_shape())

    return output_layer
```

MODEL 2 + ADAM Optimizer

Input-ReLu(512)-ReLu(128)-Relu(output) - AdamOptimizer

In [52]:

```
y_relu = mulit_layer_perceptron(x,relu_weights,biases)
cost_relu = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y
_relu, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_relu)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_relu)
```

```

cost_relu,

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([adam_optimizer,cost_relu,relu_weights],feed_dict =
{x:xs_batch, y_:ys_batch})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_relu, feed_dict = {x: data.test.images, y_: data.test.labels})
            test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu(512)-ReLu(128)-ReLu(output)-AdamOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))
            dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu(512)-ReLu(128)-ReLu(output)-AdamOptimizer")
            correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
            accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
            print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labels}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with RELU Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_21:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with RELU Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_22:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)

Output layer

Output Layer: (?, 10)

Epoch: 0001 train cost=209.421812495 test cost=210.405454614

Epoch: 0002 train cost=2.302582741 test cost=2.302565488

Epoch: 0003 train cost=2.302582741 test cost=2.302571774

Epoch: 0004 train cost=2.302582741 test cost=2.302571774

Epoch: 0005 train cost=2.302582741 test cost=2.302571774

Epoch: 0006 train cost=2.302582741 test cost=2.302571774

Epoch: 0007 train cost=2.302582741 test cost=2.302571774

Epoch: 0008 train cost=2.302582741 test cost=2.302571774

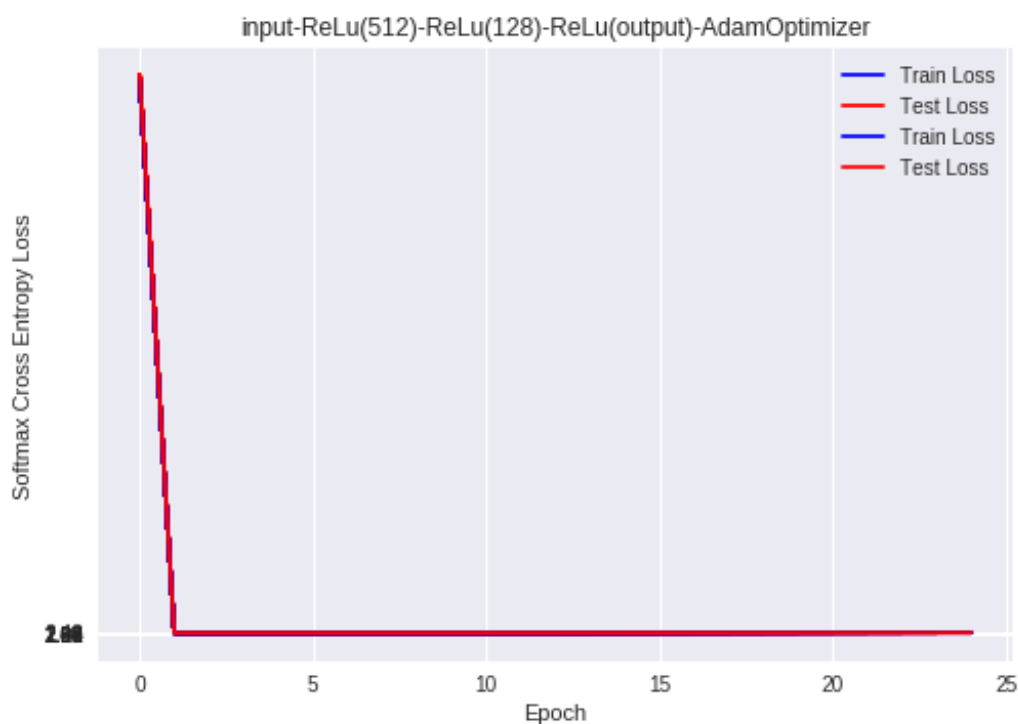
Epoch: 0009 train cost=2.302582741 test cost=2.302571774

Epoch: 0010 train cost=2.302582741 test cost=2.302571774

```

Epoch: 0011 train cost=2.302582741 test cost=2.302571774
Epoch: 0012 train cost=2.302582741 test cost=2.302571774
Epoch: 0013 train cost=2.302582741 test cost=2.302571774
Epoch: 0014 train cost=2.302582741 test cost=2.302571774
Epoch: 0015 train cost=2.302582741 test cost=2.302571774
Epoch: 0016 train cost=2.302582741 test cost=2.302571774
Epoch: 0017 train cost=2.302582741 test cost=2.302571774
Epoch: 0018 train cost=2.302582741 test cost=2.302571774
Epoch: 0019 train cost=2.302582741 test cost=2.302571774
Epoch: 0020 train cost=2.302582741 test cost=2.302571774
Epoch: 0021 train cost=2.302582741 test cost=2.302571774
Epoch: 0022 train cost=2.302582741 test cost=2.302571774
Epoch: 0023 train cost=2.302582741 test cost=2.302571774
Epoch: 0024 train cost=2.302582741 test cost=2.302571774
Epoch: 0025 train cost=2.302582741 test cost=2.302571774
Accuracy: 0.0959

```



In [53]:

```

import seaborn as s

h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")


mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")

```

```
mp.title('Weight matrix ',
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
violin_data = remove_na(group_data)
```



## MODEL 2 + GRADIENT DESCENT OPTIMIZER

In [54]:

```
y_relu = mulit_layer_perceptron(x,relu_weights,biases)
cost_relu = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y
_relu, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_relu)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_relu)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
```

```

xs_batch, ys_batch = data.train.next_batch(batch_size)
_, c, w = sess.run([sgd_optimizer, cost_relu, relu_weights], feed_dict = {
x:xs_batch, y_:ys_batch})
train_avg_cost += c/number_of_batches
c = sess.run(cost_relu, feed_dict = {x: data.test.images, y_: data.test.labels})
test_avg_cost += c/number_of_batches
epchs.append(epoch)
ytrainloss.append(train_avg_cost)
ytestloss.append(test_avg_cost)
dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu(512)-ReLu(128)-ReLu(output)-Gradient Descent Optimizer"
)

if epoch%display_step == 0:
    print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))
    dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu(512)-ReLu(128)-ReLu(output)-Gradient Descent Optimizer"
)
    correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labels}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with RELU Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_21:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with RELU Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_22:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)

Output layer

Output Layer: (?, 10)

Epoch: 0001 train cost=4.980333456 test cost=4.963308534

Epoch: 0002 train cost=1.904925619 test cost=1.912343695

Epoch: 0003 train cost=1.913920914 test cost=1.919963113

Epoch: 0004 train cost=1.906759251 test cost=1.914553289

Epoch: 0005 train cost=1.901835554 test cost=1.910970788

Epoch: 0006 train cost=1.902117114 test cost=1.909050426

Epoch: 0007 train cost=1.900284037 test cost=1.907728091

Epoch: 0008 train cost=1.898658954 test cost=1.906581226

Epoch: 0009 train cost=1.896835147 test cost=1.906401160

Epoch: 0010 train cost=1.896392504 test cost=1.904773519

Epoch: 0011 train cost=1.895304086 test cost=1.904421603

Epoch: 0012 train cost=1.894359940 test cost=1.903752665

Epoch: 0013 train cost=1.895284351 test cost=1.903683693

Epoch: 0014 train cost=1.895441476 test cost=1.903686699

Epoch: 0015 train cost=1.892990817 test cost=1.902676901

Epoch: 0016 train cost=1.893376431 test cost=1.902505441

Epoch: 0017 train cost=1.892992397 test cost=1.902061113

Epoch: 0018 train cost=1.891415631 test cost=1.901908773

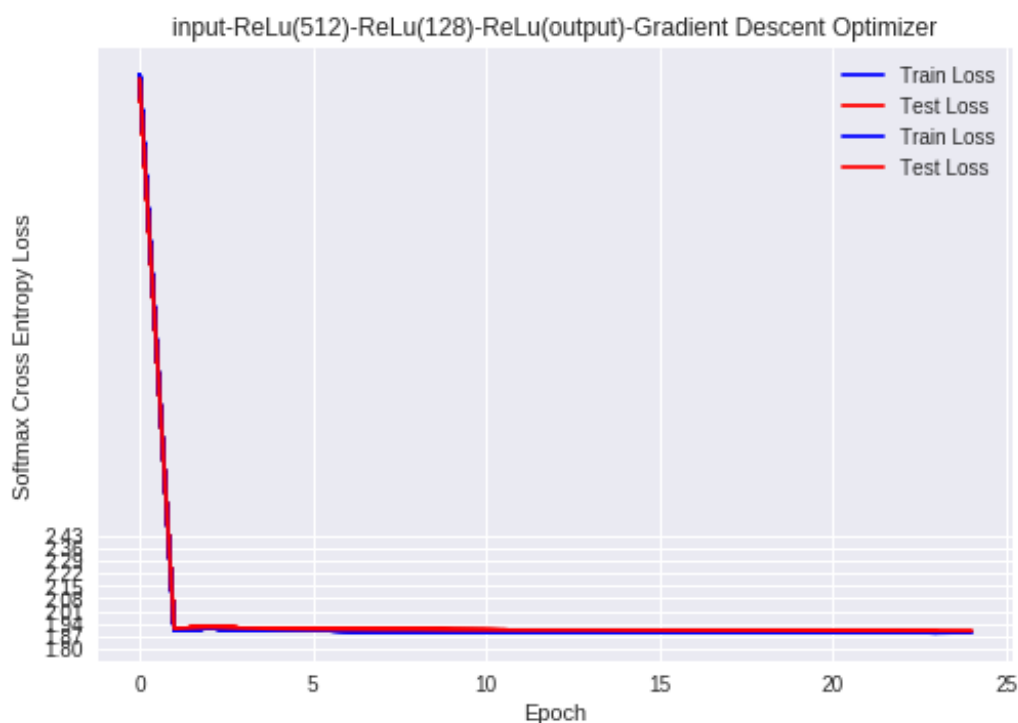
Epoch: 0019 train cost=1.892247939 test cost=1.901883442

Epoch: 0020 train cost=1.891918295 test cost=1.901420602

Epoch: 0021 train cost=1.890442915 test cost=1.900770434

Epoch: 0022 train cost=1.891120323 test cost=1.900679348

Epoch: 0023 train cost=1.893256543 test cost=1.900676042  
 Epoch: 0024 train cost=1.888824339 test cost=1.900287646  
 Epoch: 0025 train cost=1.891013288 test cost=1.899846888  
 Accuracy: 0.0954



In [55]:

```
import seaborn as s

h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

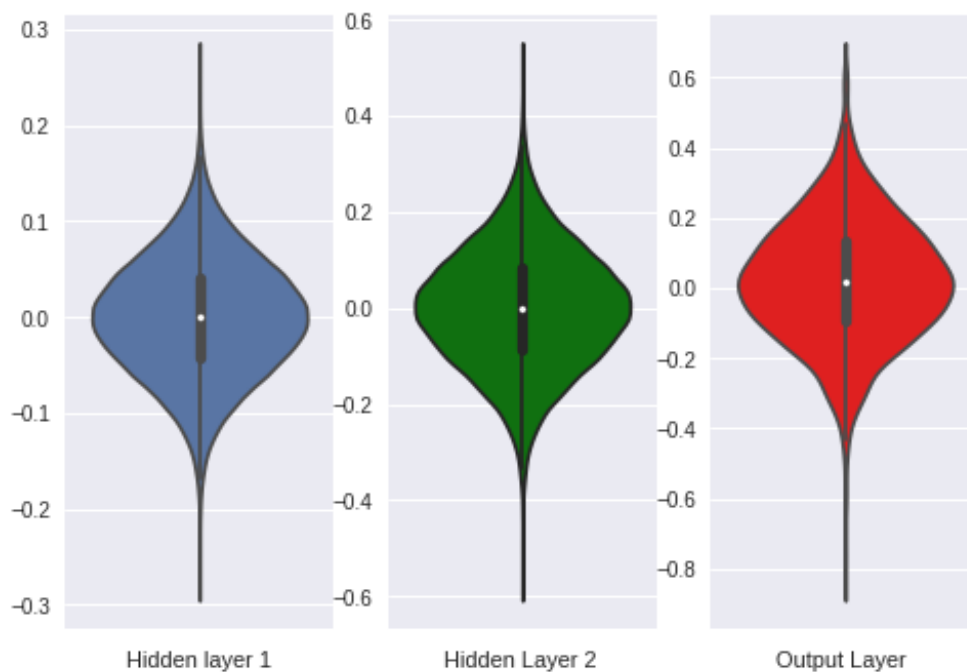
```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

Weight matrix

Weight matrix

Weight matrix





MODEL 3: INPUT - SIGMOID(BatchNormalization(512)) -  
SIGMOID(BatchNormalization(128)) - SIGMOID(output)

-> The activation functions are SIGMOID

In [0]:

```
eplpsilon = 0.01
def mulit_layer_perceptron(x,weights,biases):
    print('Input layer')
    print("X:", x.get_shape(), "W[h1]:", weights['h1'].get_shape(), 'b[h1]:',
    biases['b1'].get_shape())

    print('Hidden layer 1 with Sigmoid Activation')
    layer_1 = tf.add(tf.matmul(x,weights['h1']),biases['b1'])
    mean_1, var_1 = tf.nn.moments(layer_1,[0])
    scale_1 = tf.Variable(tf.ones([n_hidden_1]))
    beta_1 = tf.Variable(tf.zeros([n_hidden_1]))
    layer_1 = tf.nn.batch_normalization(layer_1,mean_1, var_1, scale_1,
    beta_1,eplpsilon)
    layer_1 = tf.nn.sigmoid(layer_1)
    print("Layer_1:", layer_1.get_shape(), 'W[h2]:', weights['h2'].get_shape,
    'b[b1]:', biases['b2'].get_shape())

    print("Hidden layer 2 with Sigmoid Activation")
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    mean_2, var_2 = tf.nn.moments(layer_2,[0])
    scale_2 = tf.Variable(tf.ones([n_hidden_2]))
    beta_2 = tf.Variable(tf.zeros([n_hidden_2]))
    layer_2 = tf.nn.batch_normalization(layer_2,mean_2, var_2, scale_2,
    beta_2,eplpsilon)
    layer_2 = tf.nn.sigmoid(layer_2)
    print("Layer_2:", layer_2.get_shape(), 'W[out]:', weights['out'].get_shape,
    'b3:', biases['out'].get_shape())

    print("Output layer")
    output_layer = tf.matmul(layer_2,weights['out']) + biases['out']
```

```

output_layer = tf.nn.sigmoid(output_layer)
print('Output Layer:', output_layer.get_shape())

return output_layer

```

In [0]:

### MODEL 3 + ADAM OPTIMIZER

In [69]:

```

y_sgd = mulit_layer_perceptron(x,sgd_weights,biases)
cost_sgd = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_
sgd, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_sgd)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_sgd)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([adam_optimizer,cost_sgd,sgd_weights],feed_dict = {x:
xs_batch, y_:ys_batch})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_sgd, feed_dict = {x: data.test.images, y_: data.test
.labels})
            test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid((BN)512)-sigmoid((BN)128)-sigmoid(output)-AdamOptimi
zer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_a
vg_cost), "test cost={:.9f}".format(test_avg_cost))
            dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid((BN)512)-sigmoid((BN)128)-sigmoid(output)-AdamOptimi
zer")
            correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
            accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

```

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labels}))
```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with Sigmoid Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_27:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with Sigmoid Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_28:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)

Output layer

Output Layer: (?, 10)

Epoch: 0001 train cost=2.349449171 test cost=2.352906570

Epoch: 0002 train cost=2.342291914 test cost=2.345360127

Epoch: 0003 train cost=2.266466056 test cost=2.261591890

Epoch: 0004 train cost=1.829162110 test cost=1.825918247

Epoch: 0005 train cost=1.734419727 test cost=1.733736868

Epoch: 0006 train cost=1.710431199 test cost=1.710253932

Epoch: 0007 train cost=1.706800042 test cost=1.706654499

Epoch: 0008 train cost=1.699189143 test cost=1.704543404

Epoch: 0009 train cost=1.700946758 test cost=1.704560947

Epoch: 0010 train cost=1.703178620 test cost=1.703930261

Epoch: 0011 train cost=1.696222058 test cost=1.702921508

Epoch: 0012 train cost=1.696756903 test cost=1.701534937

Epoch: 0013 train cost=1.699936199 test cost=1.701771749

Epoch: 0014 train cost=1.698799385 test cost=1.701496204

Epoch: 0015 train cost=1.695097373 test cost=1.700269287

Epoch: 0016 train cost=1.697288717 test cost=1.701145328

Epoch: 0017 train cost=1.698095374 test cost=1.701825302

Epoch: 0018 train cost=1.695608083 test cost=1.699904903

Epoch: 0019 train cost=1.695946338 test cost=1.700065361

Epoch: 0020 train cost=1.693136610 test cost=1.699424575

Epoch: 0021 train cost=1.701571848 test cost=1.703001423


Epoch: 0022 train cost=1.692584200 test cost=1.699309936


Epoch: 0023 train cost=1.696456705 test cost=1.700631369

Epoch: 0024 train cost=1.695445267 test cost=1.699066962

Epoch: 0025 train cost=1.693661974 test cost=1.700077118

Accuracy: 0.7294





In [70]:

```
import seaborn as s

h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, cicolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

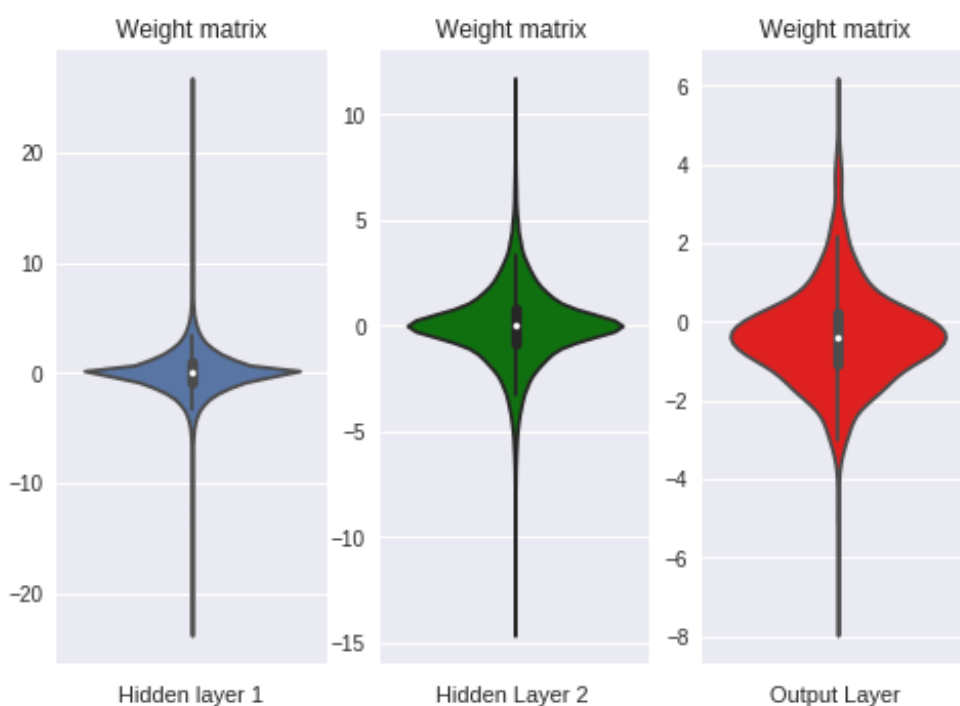
mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

```
kde_data = remove_na(group_data)
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

```
violin_data = remove_na(group_data)
```



In [71]:

```

y_sgd = mulit_layer_perceptron(x,sgd_weights,biases)
cost_sgd = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_
sgd, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_sgd)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_sgd)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([sgd_optimizer,cost_sgd,sgd_weights],feed_dict = {x:x
s_batch, y_:ys_batch})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_sgd, feed_dict = {x: data.test.images, y_: data.test
.labels})
            test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid((BN) 512)-sigmoid((BN) 128)-sigmoid(output)-AdamOptimi
zer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_a
vg_cost), "test cost={:.9f}".format(test_avg_cost))
            dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-sigmoid((BN) 512)-sigmoid((BN) 128)-sigmoid(output)-AdamOptimi
zer")
            correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
            accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
            print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labe
ls}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with Sigmoid Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_27:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

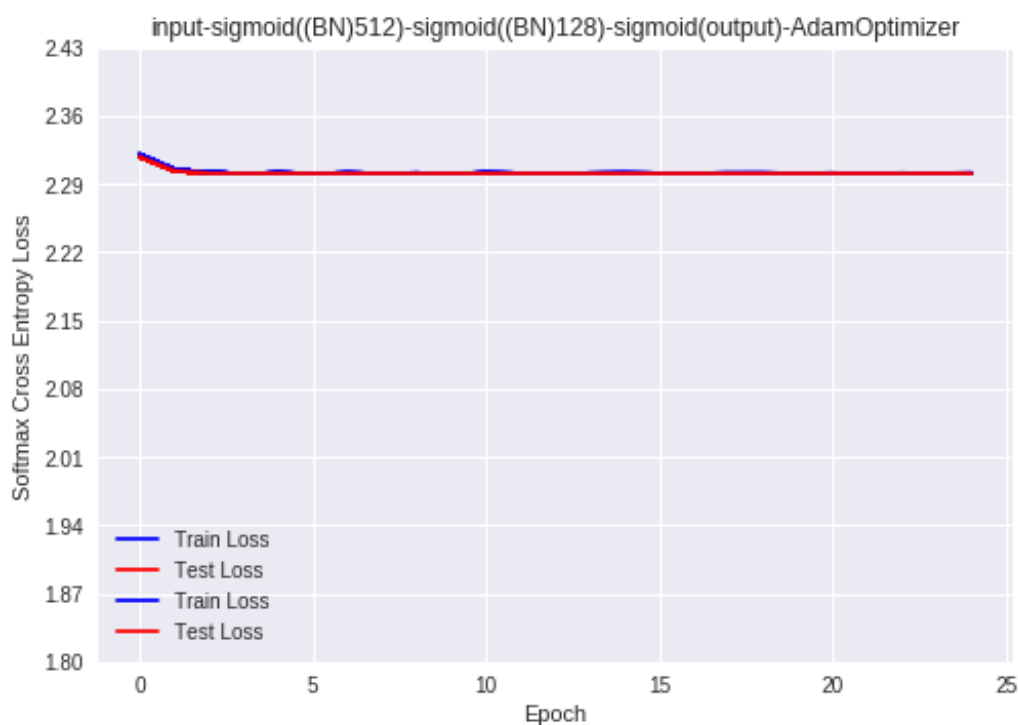
Hidden layer 2 with Sigmoid Activation

```
Layer_2: (?, 128) W[out]: <bound method Variable.get_shape of <tf.Variable  
'Variable_28:0' shape=(128, 10) dtype=float32_ref>> b3: (10,)
```

Output layer

Output Layer: (?, 10)

```
Epoch: 0001 train cost=2.320533943 test cost=2.318185780  
Epoch: 0002 train cost=2.305083630 test cost=2.303832847  
Epoch: 0003 train cost=2.301661552 test cost=2.301160088  
Epoch: 0004 train cost=2.301111967 test cost=2.301026752  
Epoch: 0005 train cost=2.301475495 test cost=2.301105400  
Epoch: 0006 train cost=2.301193979 test cost=2.301031264  
Epoch: 0007 train cost=2.301438288 test cost=2.301095143  
Epoch: 0008 train cost=2.301250159 test cost=2.301060295  
Epoch: 0009 train cost=2.301365215 test cost=2.301036874  
Epoch: 0010 train cost=2.301085242 test cost=2.301097241  
Epoch: 0011 train cost=2.301509380 test cost=2.301023938  
Epoch: 0012 train cost=2.301304674 test cost=2.301086452  
Epoch: 0013 train cost=2.301106496 test cost=2.301053021  
Epoch: 0014 train cost=2.301367014 test cost=2.301043298  
Epoch: 0015 train cost=2.301469621 test cost=2.301071037  
Epoch: 0016 train cost=2.301324597 test cost=2.301045392  
Epoch: 0017 train cost=2.301253176 test cost=2.301060603  
Epoch: 0018 train cost=2.301400961 test cost=2.301052159  
Epoch: 0019 train cost=2.301417836 test cost=2.301041339  
Epoch: 0020 train cost=2.301236304 test cost=2.301039189  
Epoch: 0021 train cost=2.301380183 test cost=2.301083261  
Epoch: 0022 train cost=2.301128041 test cost=2.301049540  
Epoch: 0023 train cost=2.301390479 test cost=2.301055999  
Epoch: 0024 train cost=2.301089820 test cost=2.301090657  
Epoch: 0025 train cost=2.301667773 test cost=2.301065579  
Accuracy: 0.1135
```



In [72]:

```
import seaborn as s
```

```
h1_w = w['h1'].flatten().reshape(-1,1)
```

```
h2_w = w['h2'].flatten().reshape(-1,1)
```

```

out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

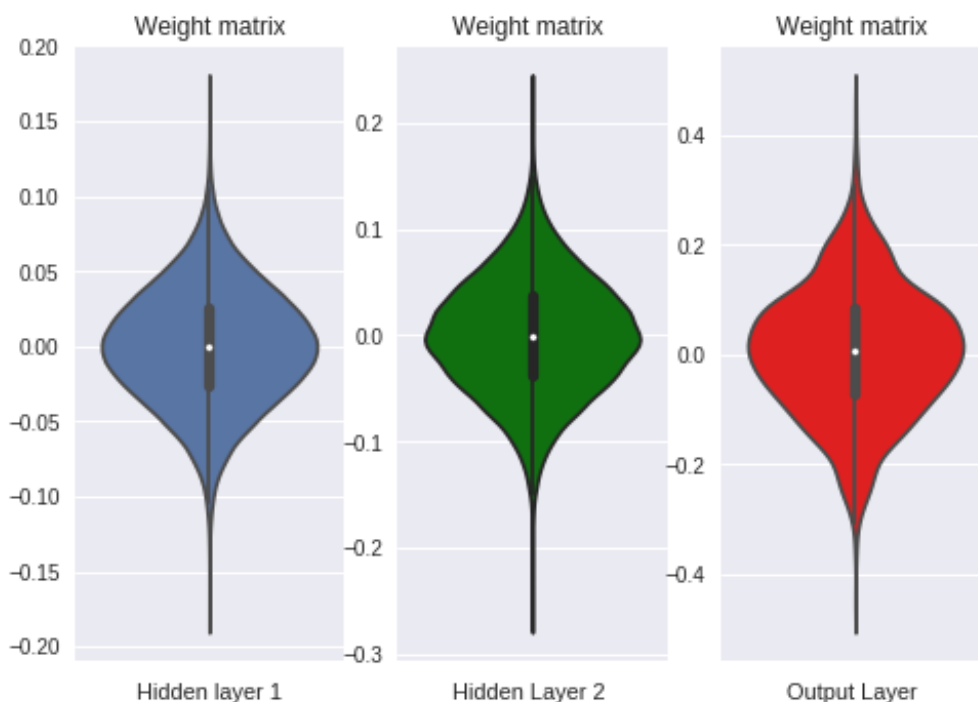
mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  violin_data = remove_na(group_data)

```



MODEL 4: INPUT - ReLu(Drop out(512)) - ReLu(Drop out(128)) - ReLu(output)

-> The activation functions are ReLU



In [0]:

```
def mulit_laver perceptron(x.weights.biases):
```

```

def multilayer_perceptron(x, weights, biases):
    print('Input layer')
    print("X:", x.get_shape(), "W[h1]:", weights['h1'].get_shape(), 'b[h1]:',
          biases['b1'].get_shape())

    print('Hidden layer 1 with RELU Activation')
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    layer_1 = tf.nn.relu(layer_1)
    print("Layer_1:", layer_1.get_shape(), 'W[h2]:', weights['h2'].get_shape(),
          'b[b1]:', biases['b2'].get_shape())
    layer_1_drop = tf.nn.dropout(layer_1, keep_prob)

    print("Hidden layer 2 with RELU Activation")
    layer_2 = tf.add(tf.matmul(layer_1_drop, weights['h2']), biases['b2'])
    layer_2 = tf.nn.relu(layer_2)
    print("Layer_2:", layer_2.get_shape(), 'W[out]:', weights['out'].get_shape(),
          'b3:', biases['out'].get_shape())
    layer_2_drop = tf.nn.dropout(layer_2, keep_prob)

    print("Output layer")
    output_layer = tf.matmul(layer_2_drop, weights['out']) + biases['out']
    output_layer = tf.nn.relu(output_layer)
    print('Output Layer:', output_layer.get_shape())

    return output_layer

```

#### MODEL 4 + ADAM OPTIMIZER

In [77]:

```

y_relu = multilayer_perceptron(x, relu_weights, biases)
cost_relu = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_relu, labels = y))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost_relu)
sgd_optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost_relu)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig, ax = mp.subplots(1, 1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _, c, w = sess.run([adam_optimizer, cost_relu, relu_weights], feed_dict =
                               {x:xs_batch, y:ys_batch, keep_prob:0.5})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_relu, feed_dict = {x: data.test.images, y: data.test

```



```

t.labels, keep_prob:0.5})
    test_avg_cost += c/number_of_batches
    epchs.append(epoch)
    ytrainloss.append(train_avg_cost)
    ytestloss.append(test_avg_cost)
    dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu(512)-ReLu(128)-ReLu(output)-AdamOptimizer")

    if epoch%display_step == 0:
        print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_a
vg_cost), "test cost={:.9f}".format(test_avg_cost))
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu((DO)512)-ReLu((DO)128)-ReLu(output)-AdamOptimizer")
        correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labe
ls, keep_prob:0.5}))

```


```

Input layer
X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)
Hidden layer 1 with RELU Activation
Layer_1: (?, 512) W[h2]: <bound method Variable.get_shape of <tf.Variable '
Variable_30:0' shape=(512, 128) dtype=float32_ref>> b[b1]: (128,)
Hidden layer 2 with RELU Activation
Layer_2: (?, 128) W[out]: <bound method Variable.get_shape of <tf.Variable
'Variable_31:0' shape=(128, 10) dtype=float32_ref>> b3: (10,)
Output layer
Output Layer: (?, 10)
Epoch: 0001 train cost=4.767311148 test cost=4.649149166
Epoch: 0002 train cost=2.302582741 test cost=2.302571774
Epoch: 0003 train cost=2.302582741 test cost=2.303704626
Epoch: 0004 train cost=2.302582741 test cost=2.305074540
Epoch: 0005 train cost=2.302582741 test cost=2.316208046
Epoch: 0006 train cost=2.302582741 test cost=2.310911512
Epoch: 0007 train cost=2.302582741 test cost=2.305555123
Epoch: 0008 train cost=2.302582741 test cost=2.302571774
Epoch: 0009 train cost=2.302582741 test cost=2.304625555
Epoch: 0010 train cost=2.302582741 test cost=2.302571774
Epoch: 0011 train cost=2.312040814 test cost=2.302571774
Epoch: 0012 train cost=2.302582741 test cost=2.302571774
Epoch: 0013 train cost=2.302582741 test cost=2.302571774
Epoch: 0014 train cost=2.302582741 test cost=2.302571774
Epoch: 0015 train cost=2.302582741 test cost=2.302571774
Epoch: 0016 train cost=2.302582741 test cost=2.302571774
Epoch: 0017 train cost=2.302582741 test cost=2.302571774
Epoch: 0018 train cost=2.302582741 test cost=2.302571774
Epoch: 0019 train cost=2.302582741 test cost=2.302571774
Epoch: 0020 train cost=2.302582741 test cost=2.302571774
Epoch: 0021 train cost=2.302582741 test cost=2.302571774
Epoch: 0022 train cost=2.302582741 test cost=2.302571774
Epoch: 0023 train cost=2.302582741 test cost=2.302571774
Epoch: 0024 train cost=2.302582741 test cost=2.302571774
Epoch: 0025 train cost=2.302582741 test cost=2.302571774
Accuracy: 0.0892

```

input-ReLu((DO)512)-ReLu((DO)128)-ReLu(output)-AdamOptimizer





In [78]:

```
import seaborn as s


h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

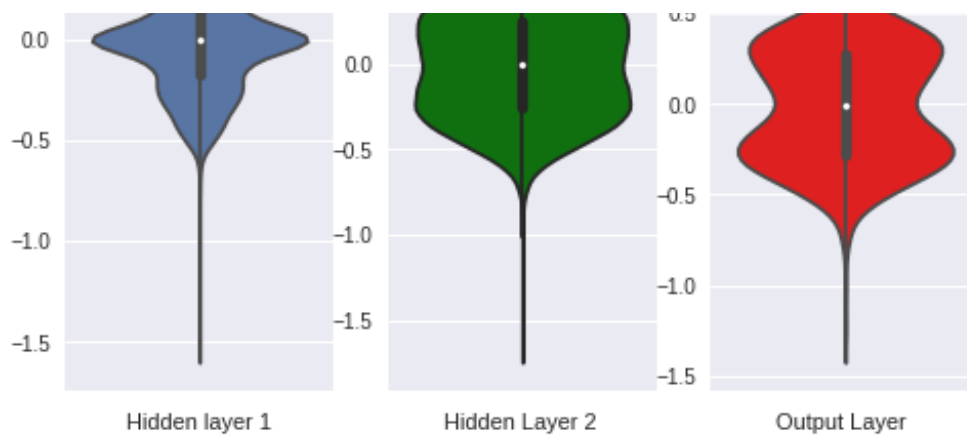
fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
  violin_data = remove_na(group_data)
```





MODEL 4 + GRADIENT DESCENT OPTIMIZER

In [81]:

```
y_relu = mulit_layer_perceptron(x,relu_weights,biases)
cost_relu = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y
_relu, labels = y_))
adam_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimi
ze(cost_relu)
sgd_optimizer =
tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cos
t_relu)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = mp.subplots(1,1)
    ax.set_xlabel('Epoch')
    ax.set_ylabel("Softmax Cross Entropy Loss")
    epchs, ytrainloss, ytestloss = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0
        test_avg_cost = 0
        number_of_batches = int(data.train.num_examples/batch_size)

        for i in range(number_of_batches):
            xs_batch, ys_batch = data.train.next_batch(batch_size)
            _,c,w = sess.run([sgd_optimizer,cost_relu,relu_weights],feed_dict = {
x:xs_batch, y_:ys_batch, keep_prob:0.5})
            train_avg_cost += c/number_of_batches
            c = sess.run(cost_relu, feed_dict = {x: data.test.images, y_: data.tes
t.labels, keep_prob:0.5})
            test_avg_cost += c/number_of_batches
        epchs.append(epoch)
        ytrainloss.append(train_avg_cost)
        ytestloss.append(test_avg_cost)
        dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu ((DO) 512)-ReLu ((DO) 128)-ReLu(output)-Gradient Descent Op
timizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".format(train_a
vg_cost), "test cost={:.9f}".format(test_avg_cost))
            dynamic_plot(epchs, ytrainloss, ytestloss, ax, np.arange(1.8, 2.5, step
=0.07), "input-ReLu ((DO) 512)-ReLu ((DO) 128)-ReLu(output)Gradient Descent Opt
... ..")
```

```

lmizer")
    correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print("Accuracy:", accuracy.eval({x: data.test.images, y_: data.test.labels, keep_prob:0.5}))

```

Input layer

X: (?, 784) W[h1]: (784, 512) b[h1]: (512,)

Hidden layer 1 with RELU Activation

Layer\_1: (?, 512) W[h2]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_30:0' shape=(512, 128) dtype=float32\_ref>> b[b1]: (128,)

Hidden layer 2 with RELU Activation

Layer\_2: (?, 128) W[out]: <bound method Variable.get\_shape of <tf.Variable 'Variable\_31:0' shape=(128, 10) dtype=float32\_ref>> b3: (10,)


Output layer


Output Layer: (?, 10)

```

Epoch: 0001 train cost=2.572081514 test cost=2.416359403
Epoch: 0002 train cost=2.304764600 test cost=2.305176154
Epoch: 0003 train cost=2.303997738 test cost=2.303121476
Epoch: 0004 train cost=2.295936758 test cost=2.296186772
Epoch: 0005 train cost=2.242655997 test cost=2.247438860
Epoch: 0006 train cost=2.196123284 test cost=2.196128598
Epoch: 0007 train cost=2.168153858 test cost=2.172993187
Epoch: 0008 train cost=2.155915421 test cost=2.158906438
Epoch: 0009 train cost=2.148236782 test cost=2.150557015
Epoch: 0010 train cost=2.137970213 test cost=2.142845969
Epoch: 0011 train cost=2.134381212 test cost=2.138372343
Epoch: 0012 train cost=2.131681863 test cost=2.133029691
Epoch: 0013 train cost=2.129146788 test cost=2.129725161
Epoch: 0014 train cost=2.120631886 test cost=2.125522679
Epoch: 0015 train cost=2.112758988 test cost=2.114596488
Epoch: 0016 train cost=2.074142833 test cost=2.074843818
Epoch: 0017 train cost=2.000587242 test cost=1.993752909
Epoch: 0018 train cost=1.928366858 test cost=1.918500996
Epoch: 0019 train cost=1.871250831 test cost=1.853919181
Epoch: 0020 train cost=1.698395027 test cost=1.668408006
Epoch: 0021 train cost=1.550084251 test cost=1.521005119
Epoch: 0022 train cost=1.475075945 test cost=1.458681367
Epoch: 0023 train cost=1.454468053 test cost=1.427334798
Epoch: 0024 train cost=1.423312850 test cost=1.406613092
Epoch: 0025 train cost=1.407162792 test cost=1.387703720
Accuracy: 0.1028

```





In [82]:

```
import seaborn as s


h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = mp.figure()
mp.subplot(1,3,1)
mp.title("Weight matrix")
ax = s.violinplot(y = h1_w, clolr = 'b')
mp.xlabel("Hidden layer 1")

mp.subplot(1, 3, 2)
mp.title("Weight matrix ")
ax = s.violinplot(y=h2_w, color='g')
mp.xlabel('Hidden Layer 2 ')

mp.subplot(1, 3, 3)
mp.title("Weight matrix ")
ax = s.violinplot(y=out_w,color='r')
mp.xlabel('Output Layer ')
mp.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 kde\_data = remove\_na(group\_data)  
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 violin\_data = remove\_na(group\_data)



## CONCLUSION:

-> MODEL 1

-> SIGMOID ACTIVATION FUNCTIONS

-> ADAM OPTIMIZER

-> TRAIN LOSS - 2.416  
-> TEST LOSS - 2.414  
-> ACCURACY - 10.3%

-> GRADIENT DESCENT OPTIMIZER

-> TRAIN LOSS - 1.5983  
-> TEST LOSS - 1.5925  
-> ACCURACY - 89.36%

-> MODEL 2

-> RELU ACTIVATION FUNCTIONS

-> ADAM OPTIMIZER

-> TRAIN LOSS - 2.3625  
-> TEST LOSS - 2.3025  
-> ACCURACY - 9.59%

-> GRADIENT DESCENT OPTIMIZER

-> TRAIN LOSS - 1.8910  
-> TEST LOSS - 1.8998  
-> ACCURACY - 9.54%

-> MODEL 3

-> SIGMOID ACTIVATION FUNCTIONS WITH

BATCH NORMALIZATION

-> ADAM OPTIMIZER

1.6936  
-> TRAIN LOSS -

1.70007  
-> TEST LOSS -

-> ACCURACY - 72.9%

-> GRADIENT DESCENT OPTIMIZER

2.3016

-> TRAIN LOSS -

2.3010

-> TEST LOSS -

-> ACCURACY - 11.0

-> MODEL 4

-> RELU ACTIVATION FUNCTION WITH

ROP OUT

-> ADAM OPTIMIZER

2.3025

-> TRAIN LOSS -

2.3025

-> TEST LOSS -

8.92

-> ACCURACY -

IZER

-> GRADIENT DESCENT OPTI

1.4071

-> TRAIN LOSS -

1.3877

-> TEST LOSS -

10.28

-> ACCURACY -

